

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn as sk
```

```
data=pd.read_csv('C:\Rohith\Backup\Desktop\SEM 6\Machine Learning Lab\Practices\Diabetes\diabetes.csv')
```

```
data.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	

```
data.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	I
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

```
data.shape
```

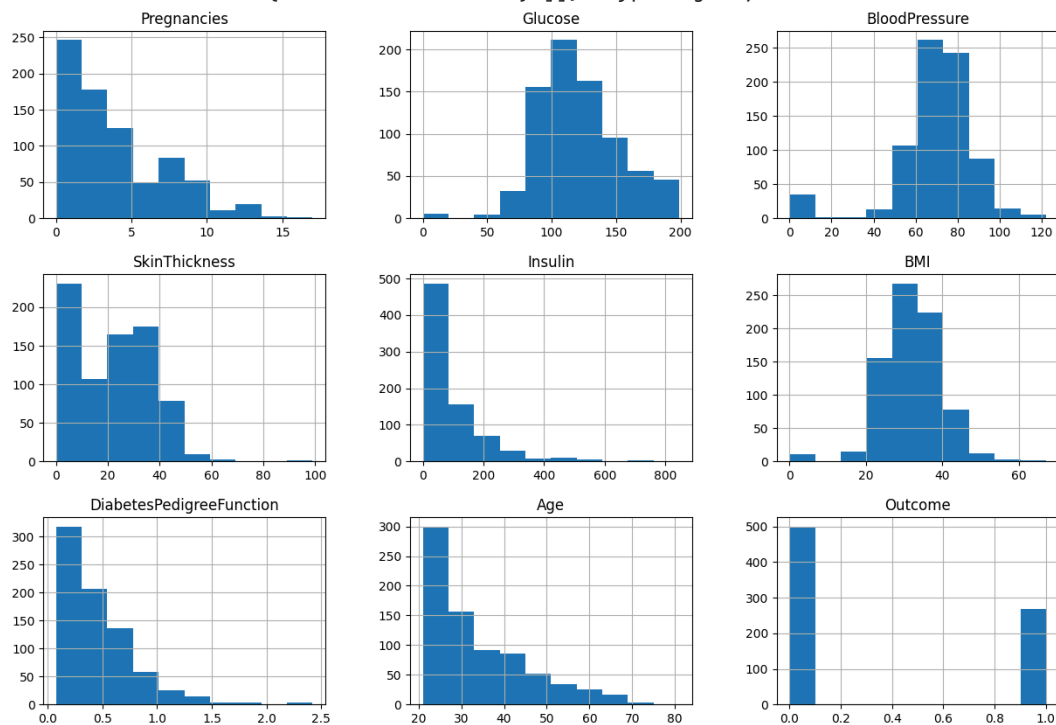
```
(768, 9)
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                        768 non-null    int64
4   Insulin                              768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
data.hist(figsize=(15,10))
```

```
array([[<Axes: title={'center': 'Pregnancies'}>,
        <Axes: title={'center': 'Glucose'}>,
        <Axes: title={'center': 'BloodPressure'}>],
       [<Axes: title={'center': 'SkinThickness'}>,
        <Axes: title={'center': 'Insulin'}>,
        <Axes: title={'center': 'BMI'}>],
       [<Axes: title={'center': 'DiabetesPedigreeFunction'}>,
        <Axes: title={'center': 'Age'}>,
        <Axes: title={'center': 'Outcome'}>]], dtype=object)
```



```
#Checking the null values
data.isnull().sum()
```


```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

```
data.duplicated().sum()
```

```
0
```

```
data.plot()
```

↩ <Axes: >



```
data['Insulin'].unique()
```


↩ array([0, 94, 168, 88, 543, 846, 175, 230, 83, 96, 235, 146, 115, 140, 110, 245, 54, 192, 207, 70, 240, 82, 36, 23, 300, 342, 304, 142, 128, 38, 100, 90, 270, 71, 125, 176, 48, 64, 228, 76, 220, 40, 152, 18, 135, 495, 37, 51, 99, 145, 225, 49, 50, 92, 325, 63, 284, 119, 204, 155, 485, 53, 114, 105, 285, 156, 78, 130, 55, 58, 160, 210, 318, 44, 190, 280, 87, 271, 129, 120, 478, 56, 32, 744, 370, 45, 194, 680, 402, 258, 375, 150, 67, 57, 116, 278, 122, 545, 75, 74, 182, 360, 215, 184, 42, 132, 148, 180, 205, 85, 231, 29, 68, 52, 255, 171, 73, 108, 43, 167, 249, 293, 66, 465, 89, 158, 84, 72, 59, 81, 196, 415, 275, 165, 579, 310, 61, 474, 170, 277, 60, 14, 95, 237, 191, 328, 250, 480, 265, 193, 79, 86, 326, 188, 106, 65, 166, 274, 77, 126, 330, 600, 185, 25, 41, 272, 321, 144, 15, 183, 91, 46, 440, 159, 540, 200, 335, 387, 22, 291, 392, 178, 127, 510, 16, 112], dtype=int64)

```
max(data['Insulin'].unique())
```

↩ 846

```
plt.boxplot(data['Insulin'])
```

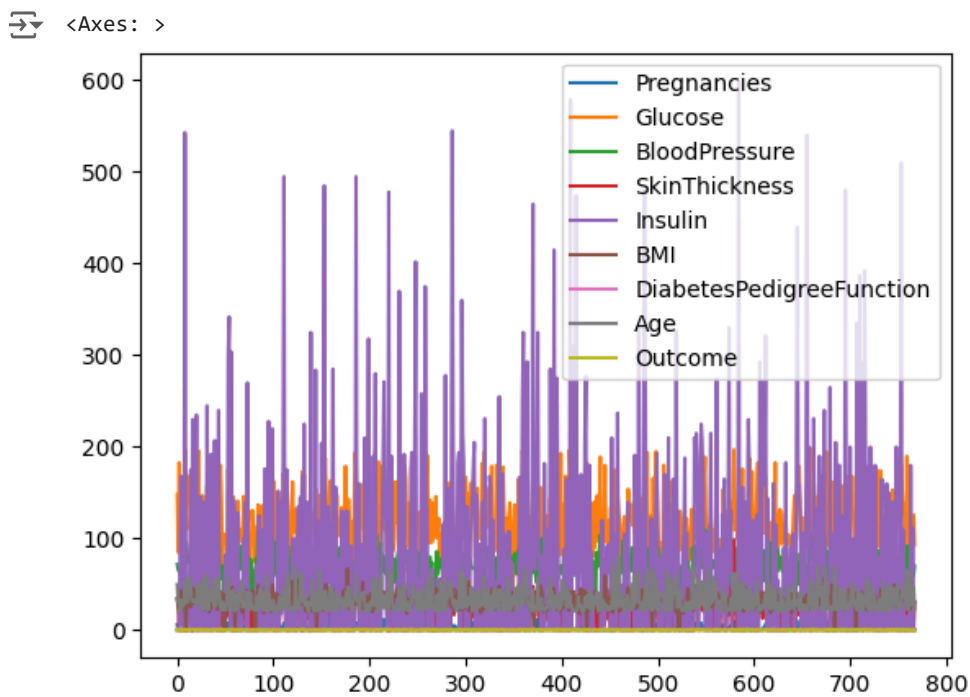
```
{'whiskers': [<matplotlib.lines.Line2D at 0x13008342980>,  
<matplotlib.lines.Line2D at 0x13008342c20>],  
'caps': [<matplotlib.lines.Line2D at 0x13008342e00>,  
<matplotlib.lines.Line2D at 0x130083430a0>],  
'boxes': [<matplotlib.lines.Line2D at 0x130083426e0>],  
'medians': [<matplotlib.lines.Line2D at 0x13008343340>],  
'fliers': [<matplotlib.lines.Line2D at 0x130083435e0>],  
'means': []}
```




```
df=data[data['Insulin']<=600]  
df
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

```
df.plot()
```



```
sns.heatmap(df.corr())
```



Training and Testing

```
X=df.drop('Outcome',axis=1)
y=df['Outcome']
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
```

```
print("X_train = ",X_train.shape)
print("X_test = ",X_test.shape)
print("y_train = ",y_train.shape)
print("y_test = ",y_test.shape)
```

```
⇒ X_train = (612, 8)
   X_test = (153, 8)
   y_train = (612,)
   y_test = (153,)
```

✓ Model 1 Linear Regression

```
from sklearn.linear_model import LinearRegression
li=LinearRegression()
li
```

```
⇒ ▾ LinearRegression ⓘ ?
   LinearRegression()
```

```
li.fit(X_train,y_train)
y_pred=li.predict(X_test)
accuracy=li.score(X_test,y_test)
accuracy
```

```
⇒ 0.2775862185270629
```

```
from sklearn.metrics import mean_absolute_error,mean_squared_error
```

```
# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_pred)
print("Mean Absolute Error (MAE):", mae)
```

```
# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error (MSE):", mse)
```

```
# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print("Root Mean Squared Error (RMSE):", rmse)
```


```
⇒ Mean Absolute Error (MAE): 0.3199159992882768
   Mean Squared Error (MSE): 0.15374708271597134
   Root Mean Squared Error (RMSE): 0.3921059585315828
```

```
#Plotting ROC curve
from sklearn.metrics import roc_curve, auc, accuracy_score

fpr, tpr, threshold = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()
```

➡ No artists with labels found to put in legend. Note that artists whose label start w




```
from sklearn.metrics import confusion_matrix

# Convert regression output to binary classification
threshold = y_test.mean() # Example threshold, you can set your own
y_test_class = (y_test > threshold).astype(int)
y_pred_class = (y_pred > threshold).astype(int)

# Compute confusion matrix
cm = confusion_matrix(y_test_class, y_pred_class)
accuracy = accuracy_score(y_test_class, y_pred_class)
print("Linear Regression Accuracy: \n ", accuracy)
print("Linear Regression Confusion Matrix: \n", cm)
```

➡ Linear Regression Accuracy:
0.6928104575163399
Linear Regression Confusion Matrix:
[[67 39]
[8 39]]

```
# Plot confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test_class,y_pred_class))
```


	precision	recall	f1-score	support
0	0.89	0.63	0.74	106
1	0.50	0.83	0.62	47
accuracy			0.69	153
macro avg	0.70	0.73	0.68	153
weighted avg	0.77	0.69	0.70	153

✓ Model 2 Logistic Regression

```
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
lr
```

LogisticRegression ⓘ ?
 LogisticRegression()


```
lr.fit(X_train,y_train)
y_pred=lr.predict(X_test)
accuracy=lr.score(X_test,y_test)
accuracy
```

 c:\Python310\lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression


```
n_iter_i = _check_optimize_result(
0.8300653594771242
```

```
#Plotting ROC curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color="darkorange",lw=2)
plt.plot([0,1],[0,1],color="navy",lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

 No artists with labels found to put in legend. Note that artists whose label start w



```
#Confusion matrix
from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred)
print("Logistic REgression Confusion matrix: \n",cm)
```

```


Logistic REgression Confusion matrix:
[[95 11]
 [15 32]]

```

```

#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.title('Logistic Regression Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```



```

from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))

```

```

precision    recall  f1-score   support

0           0.86      0.90      0.88       106
1           0.74      0.68      0.71        47

 accuracy          0.83       153
 macro avg          0.80       153
weighted avg          0.83       153

```

Model 3 SVM

```

from sklearn.svm import SVC
svm_classifier=SVC(probability=True)
svm_classifier

```

↗ SVC ⓘ ?
SVC(probability=True)

```
svm_classifier.fit(X_train,y_train)
y_pred=svm_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```


↗ 0.8235294117647058

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1], [0,1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

↗ No artists with labels found to put in legend. Note that artists whose label start w




```
#Confusion matrix
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print("SVM Confusion Matrix: \n",cm)
```

↗ SVM Confusion Matrix:
[[96 10]
[17 30]]

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('SVM Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

```
precision    recall  f1-score   support

0           0.85        0.91        0.88         106
1           0.75        0.64        0.69          47

 accuracy          0.82         153
 macro avg          0.80        0.77        0.78         153
weighted avg          0.82        0.82        0.82         153
```

Model 4 Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
```

```
dt_classifier=DecisionTreeClassifier()
dt_classifier
```

```
DecisionTreeClassifier
```

DecisionTreeClassifier()

```
dt_classifier.fit(X_train,y_train)
y_pred=dt_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```


↗ 0.7516339869281046

```
#Plotting ROC Curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1], [0,1], color='navy',linestyle='--',lw=2)
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve')
plt.legend(loc='lower right')
plt.show()
```

↗ No artists with labels found to put in legend. Note that artists whose label start w




```
#Confusion Matrix
from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred)
print("Decision Tree Confusion Matrix: \n",cm)
```

↗ Decision Tree Confusion Matrix:
[[80 26]
[12 35]]

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Decision Tree Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

```
precision    recall  f1-score   support

      0       0.87      0.75      0.81       106
      1       0.57      0.74      0.65        47

 accuracy          0.75       153
 macro avg       0.72      0.75      0.73       153
weighted avg       0.78      0.75      0.76       153
```

✓ Model 5 Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_classifier=RandomForestClassifier()
rf_classifier
```

```
RandomForestClassifier
```

RandomForestClassifier()

```
rf_classifier.fit(X_train,y_train)
y_pred=rf_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```

↗ 0.7908496732026143


#Plotting ROC curve

```
from sklearn.metrics import roc_curve, auc, accuracy_score
```

```
fpr, tpr, threshold = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)
```

```
plt.plot(fpr, tpr, color='darkorange', lw=2)
plt.plot([0,1], [0,1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

↗ No artists with labels found to put in legend. Note that artists whose label start w




#Confusion Matrix

```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print("Random Forest Confusion Matrix: \n",cm)
```

↗ Random Forest Confusion Matrix:
[[86 20]
[12 35]]

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Random Forest Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

```
precision    recall  f1-score   support

      0       0.88      0.81      0.84        106
      1       0.64      0.74      0.69         47

 accuracy          0.79        153
 macro avg       0.76      0.78      0.76        153
weighted avg       0.80      0.79      0.79        153
```

✓ Model 6 K Neighbors

```
from sklearn.neighbors import KNeighborsClassifier

knn_classifier=KNeighborsClassifier()
knn_classifier
```

```
KNeighborsClassifier
KNeighborsClassifier()
```



```
knn_classifier.fit(X_train,y_train)
y_pred=knn_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```


↗ 0.7581699346405228

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

↗ No artists with labels found to put in legend. Note that artists whose label start w



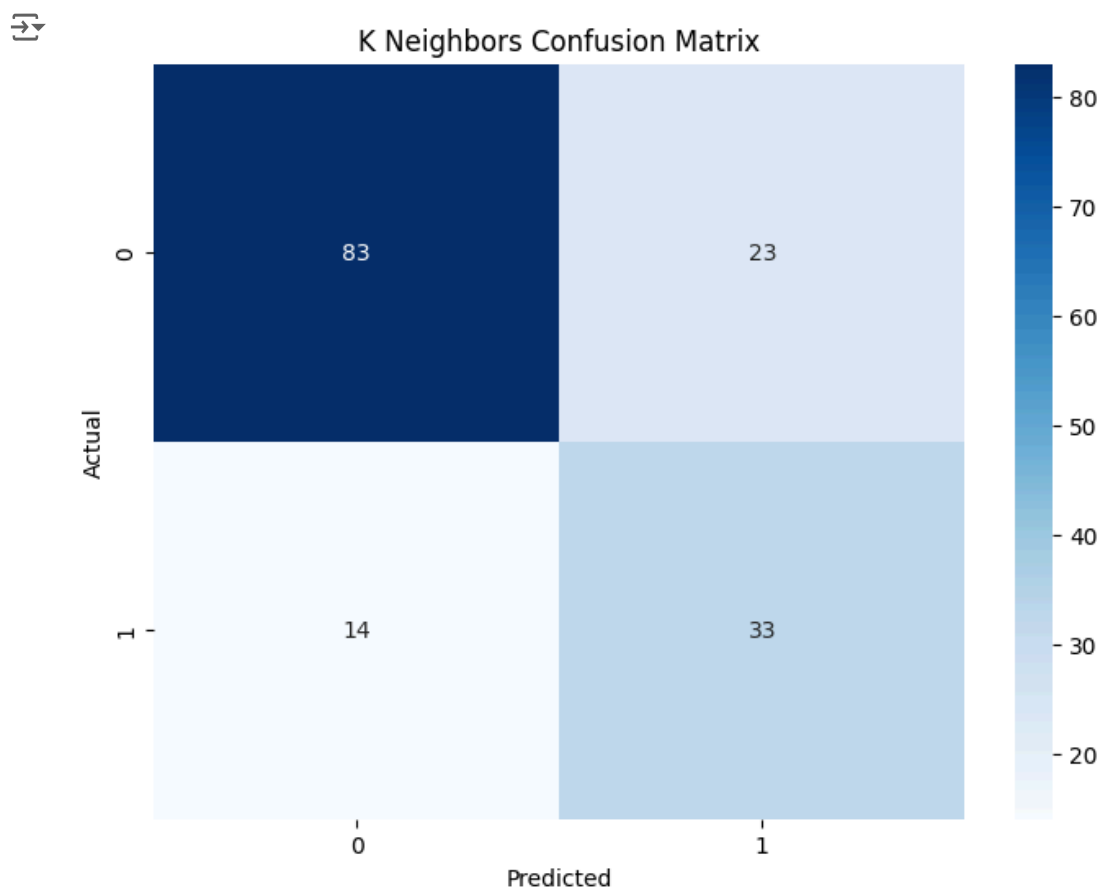
```
#Confusion Matrix
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print("K Neighbors Confusion Matrix: \n",cm)
```

↗ K Neighbors Confusion Matrix:
[[83 23]
[14 33]]

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('K Neighbors Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.78	0.82	106
1	0.59	0.70	0.64	47
accuracy			0.76	153
macro avg	0.72	0.74	0.73	153
weighted avg	0.77	0.76	0.76	153

Model 7 PLA

```
from sklearn.linear_model import Perceptron
```

```
pla_classifier=Perceptron(max_iter=1000,tol=1e-3,random_state=42)
pla_classifier
```

Perceptron ⓘ ?
Perceptron(random_state=42)

```

pla_classifier.fit(X_train,y_train)
y_pred=pla_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy

```

↗ 0.5228758169934641

```

#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score

```

```

fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)


```

```

plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()

```

↗ No artists with labels found to put in legend. Note that artists whose label start w



```

#Confusion Matrix
from sklearn.metrics import confusion_matrix

```

```

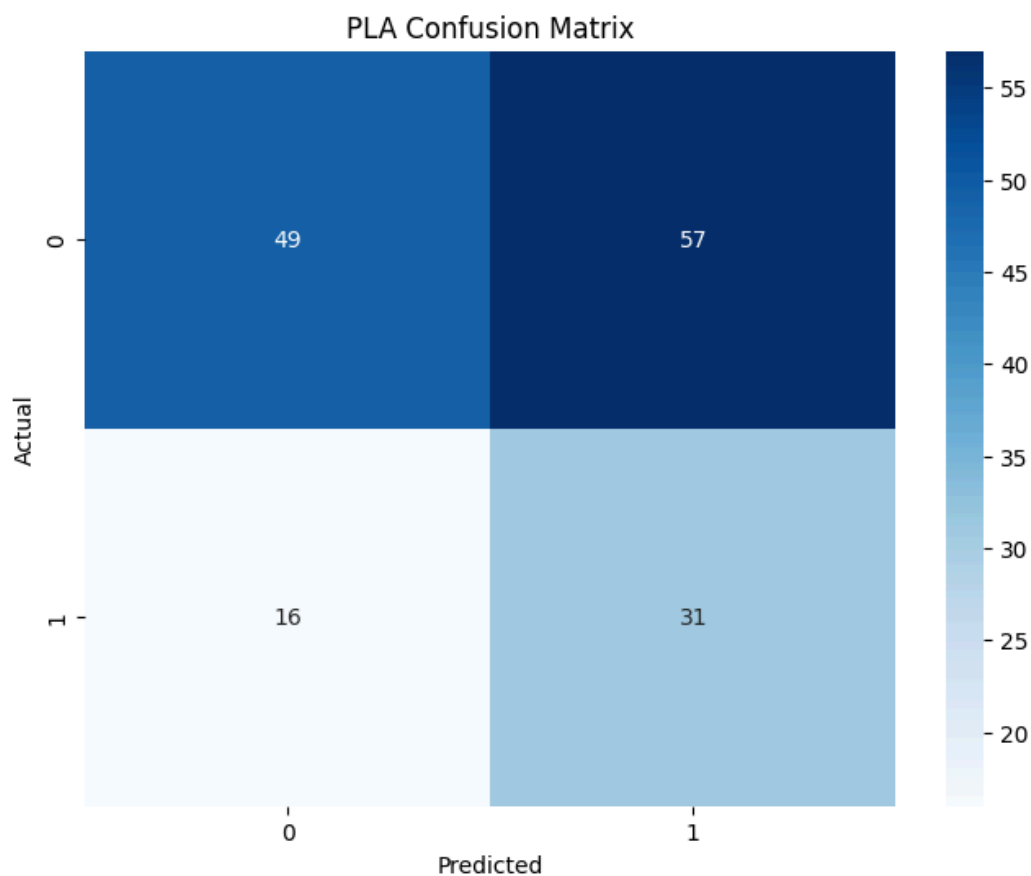
cm=confusion_matrix(y_test,y_pred)
print("PLA Matrix: \n",cm)

```

↗ PLA Matrix:
[[49 57]
[16 31]]

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('PLA Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.75	0.46	0.57	106
1	0.35	0.66	0.46	47
accuracy			0.52	153
macro avg	0.55	0.56	0.52	153
weighted avg	0.63	0.52	0.54	153

Model 8 MLP

```
from sklearn.neural_network import MLPClassifier
```

```
mlp_classifier=MLPClassifier(hidden_layer_sizes=(100,),max_iter=300,random_state=42)
mlp_classifier
```



MLPClassifier

MLPClassifier(max_iter=300, random_state=42)

```
mlp_classifier.fit(X_train,y_train)
y_pred=mlp_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```


```
0.738562091503268
```

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

```
No artists with labels found to put in legend. Note that artists whose label start w
```



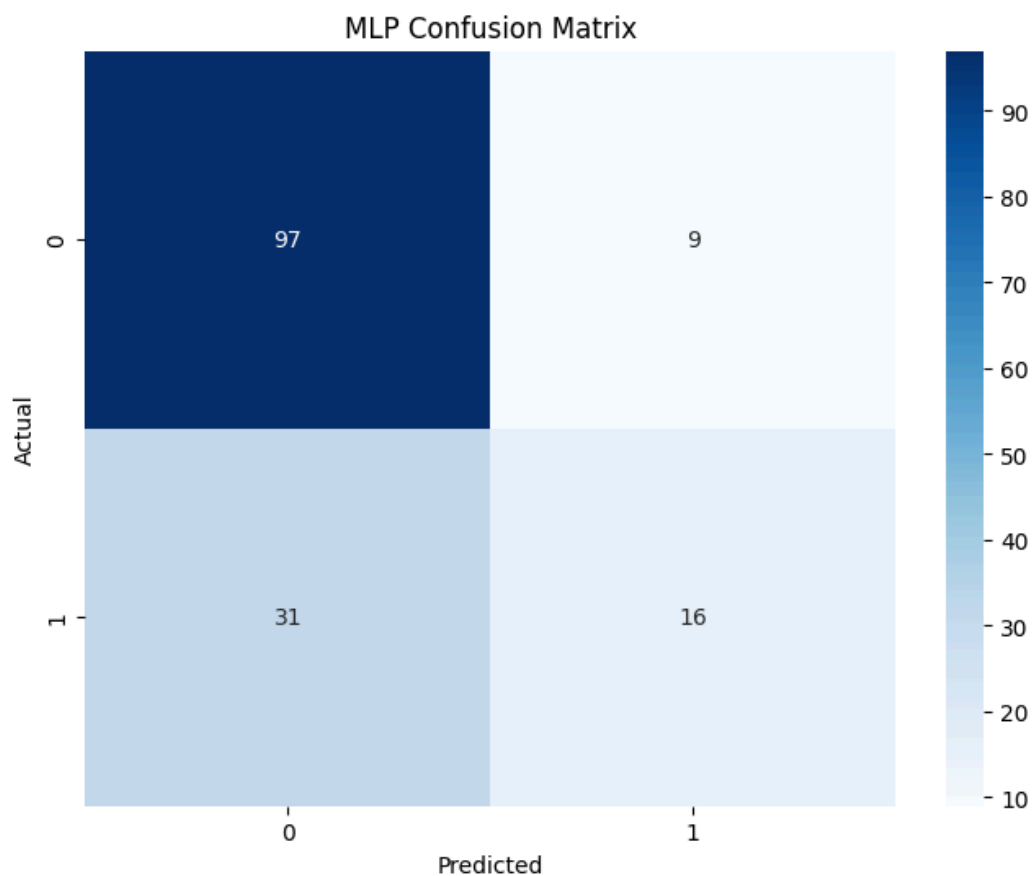
```
#Confusion Matrix
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print("MLP Matrix: \n",cm)
```

```
MLP Matrix:
[[97  9]
 [31 16]]
```

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('MLP Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.76	0.92	0.83	106
1	0.64	0.34	0.44	47
accuracy			0.74	153
macro avg	0.70	0.63	0.64	153
weighted avg	0.72	0.74	0.71	153

Model 9 Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
```

```
nb_classifier=GaussianNB()
nb_classifier
```



▼ GaussianNB ⓘ ?

GaussianNB()

```
nb_classifier.fit(X_train,y_train)
y_pred=nb_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```


```
0.7712418300653595
```

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

```
No artists with labels found to put in legend. Note that artists whose label start w
```




```
#Confusion Matrix
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print("Naive Bayes Matrix: \n",cm)
```

```
Naive Bayes Matrix:
[[88 18]
 [17 30]]
```

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Naive Bayes Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.84	0.83	0.83	106
1	0.62	0.64	0.63	47
accuracy			0.77	153
macro avg	0.73	0.73	0.73	153
weighted avg	0.77	0.77	0.77	153

Model 10 KMeans Clustering

```
from sklearn.cluster import KMeans
```

```
kmeans=KMeans(n_clusters=3,random_state=42)
kmeans
```



```
KMeans
KMeans(n_clusters=3, random_state=42)
```




```
kmeans.fit(X)
y_pred=kmeans.predict(X)

from sklearn.metrics import silhouette_score
silhouette_avg=silhouette_score(X,y_pred)
print("silhouette_score: ",silhouette_avg)
```

```
silhouette_score: 0.4989562777213992
```

```
plt.figure(figsize=(10,6))
plt.scatter(X.iloc[:,0],X.iloc[:,1],c=y_pred,s=50,cmap='viridis')
centers=kmeans.cluster_centers_
plt.scatter(centers[:,0],centers[:,1],c='red',s=200,alpha=0.75)
plt.title('KMeans Clustering')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```



✓ Model 11 Dimensionality Reduction using PCA

```
from sklearn.decomposition import PCA
pca=PCA(n_components=2)
pca
```

```
PCA
PCA(n_components=2)
```

```
X_pca=pca.fit_transform(X)
```

```
plt.figure(figsize=(8,6))
plt.scatter(X_pca[:,0],X_pca[:,1],cmap='viridis')
plt.title('Dimensionality Reduction using PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```

→ C:\Users\rohit\AppData\Local\Temp\ipykernel_17604\2795715780.py:2: UserWarning: No data found for the following column(s): Principal Component 1
plt.scatter(X_pca[:,0],X_pca[:,1],cmap='viridis')

