



0	Gender	2111	non-null	object
1	Age	2111	non-null	float64
2	Height	2111	non-null	float64
3	Weight	2111	non-null	float64
4	family_history_with_overweight	2111	non-null	object
5	FAVC	2111	non-null	object
6	FCVC	2111	non-null	float64
7	NCP	2111	non-null	float64
8	CAEC	2111	non-null	object
9	SMOKE	2111	non-null	object
10	CH2O	2111	non-null	float64
11	SCC	2111	non-null	object
12	FAF	2111	non-null	float64
13	TUE	2111	non-null	float64
14	CALC	2111	non-null	object
15	MTRANS	2111	non-null	object
16	NObeyesdad	2111	non-null	object

dtypes: float64(8), object(9)  
memory usage: 280.5+ KB

data.shape

⇒ (2111, 17)

data.isnull().sum()

⇒

Gender	0
Age	0
Height	0
Weight	0
family_history_with_overweight	0
FAVC	0
FCVC	0
NCP	0
CAEC	0
SMOKE	0
CH2O	0
SCC	0
FAF	0
TUE	0
CALC	0
MTRANS	0
NObeyesdad	0

dtype: int64

data.duplicated().sum()

⇒ 24

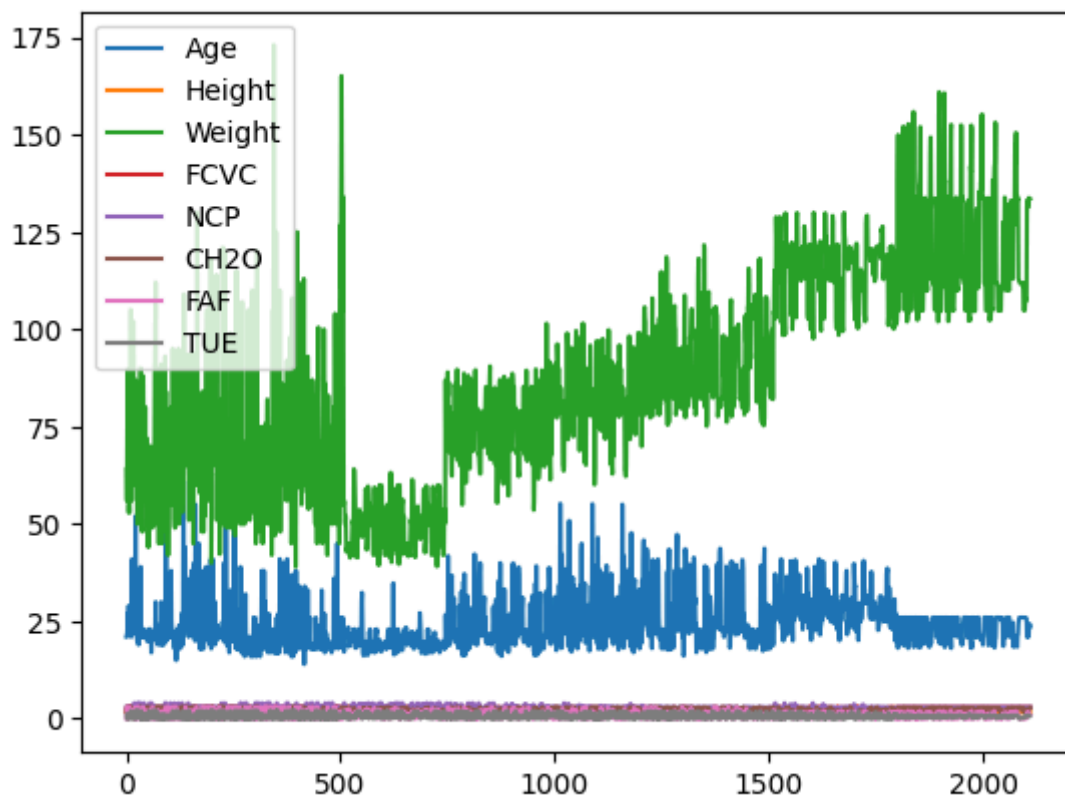
data.drop\_duplicates(inplace=True)

data.shape

⇒ (2087, 17)

```
data.plot()
```

↩ <Axes: >



```
data['Weight'].unique()
```


↩ array([ 64. , 56. , 77. , ..., 133.689352, 133.346641,  
 133.472641])

```
max(data['Weight'].unique())
```

↩ 173.0

```
plt.boxplot(data['Weight'])
```

```
➡ {'whiskers': [<matplotlib.lines.Line2D at 0x1d5a70e20b0>,  
               <matplotlib.lines.Line2D at 0x1d5a70e2230>],  
   'caps': [<matplotlib.lines.Line2D at 0x1d5a70e24d0>,  
            <matplotlib.lines.Line2D at 0x1d5a70e2770>],  
   'boxes': [<matplotlib.lines.Line2D at 0x1d5a70e1e10>],  
   'medians': [<matplotlib.lines.Line2D at 0x1d5a70e2a10>],  
   'fliers': [<matplotlib.lines.Line2D at 0x1d5a70e2cb0>],  
   'means': []}
```

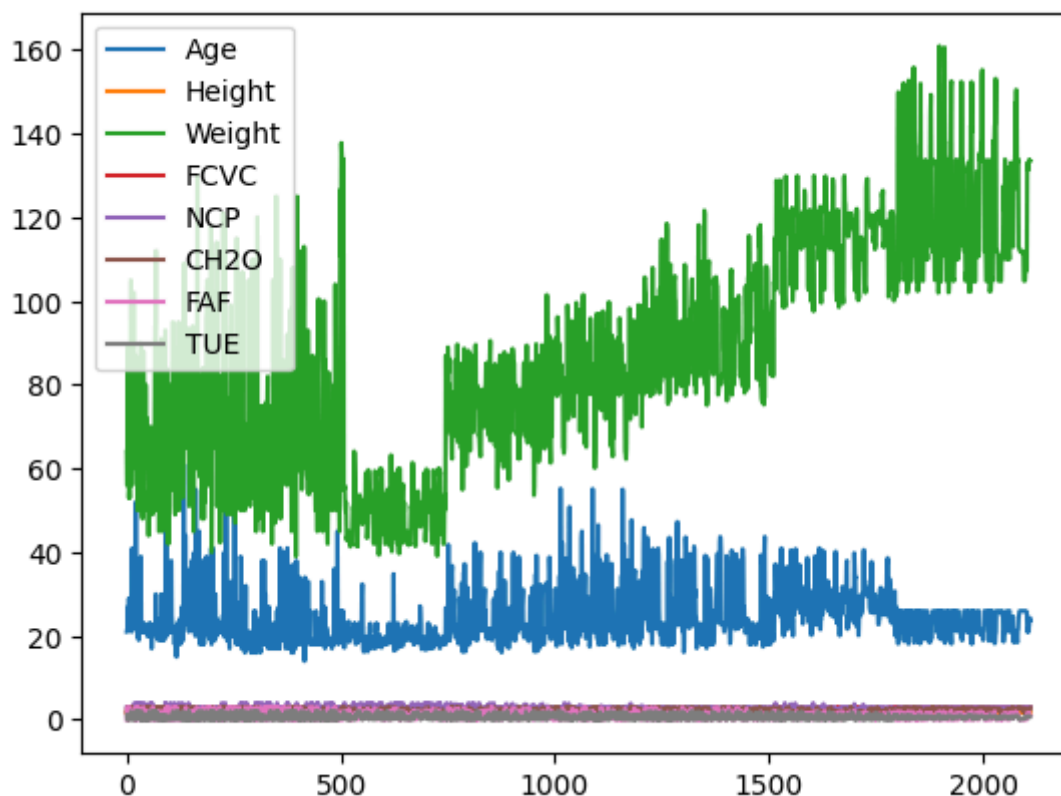


```
df=data[data['Weight']<=165]  
df.shape
```

```
➡ (2085, 17)
```

```
df.plot()
```

↔ <Axes: >



df.info()

↔ <class 'pandas.core.frame.DataFrame'>

Index: 2085 entries, 0 to 2110

Data columns (total 17 columns):

#	Column	Non-Null Count	Dtype
0	Gender	2085 non-null	object
1	Age	2085 non-null	float64
2	Height	2085 non-null	float64
3	Weight	2085 non-null	float64
4	family_history_with_overweight	2085 non-null	object
5	FAVC	2085 non-null	object
6	FCVC	2085 non-null	float64
7	NCP	2085 non-null	float64
8	CAEC	2085 non-null	object
9	SMOKE	2085 non-null	object
10	CH2O	2085 non-null	float64
11	SCC	2085 non-null	object
12	FAF	2085 non-null	float64
13	TUE	2085 non-null	float64
14	CALC	2085 non-null	object
15	MTRANS	2085 non-null	object
16	NObeyesdad	2085 non-null	object

dtypes: float64(8), object(9)

memory usage: 293.2+ KB

```
from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()
```

```

le.fit(df['Gender'])
df['Gender']=le.transform(df['Gender'])

le.fit(df['family_history_with_overweight'])
df['family_history_with_overweight']=le.transform(df['family_history_with_overweight'])

le.fit(df['FAVC'])
df['FAVC']=le.transform(df['FAVC'])

le.fit(df['CAEC'])
df['CAEC']=le.transform(df['CAEC'])

le.fit(df['SMOKE'])
df['SMOKE']=le.transform(df['SMOKE'])


le.fit(df['SCC'])
df['SCC']=le.transform(df['SCC'])

le.fit(df['CALC'])
df['CALC']=le.transform(df['CALC'])

le.fit(df['MTRANS'])
df['MTRANS']=le.transform(df['MTRANS'])

le.fit(df['NObeyesdad'])
df['NObeyesdad']=le.transform(df['NObeyesdad'])

```

 C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:2: SettingWithCopyWarn  
 A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
 df['Gender']=le.transform(df['Gender'])

C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:5: SettingWithCopyWarn  
 A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
 df['family\_history\_with\_overweight']=le.transform(df['family\_history\_with\_overweigh

C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:8: SettingWithCopyWarn  
 A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
 df['FAVC']=le.transform(df['FAVC'])

C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:11: SettingWithCopyWar  
 A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
 df['CAEC']=le.transform(df['CAEC'])

C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:14: SettingWithCopyWar  
 A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
`df['SMOKE']=le.transform(df['SMOKE'])`  
 C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:17: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
`df['SCC']=le.transform(df['SCC'])`  
 C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:20: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
`df['CALC']=le.transform(df['CALC'])`  
 C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:23: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
`df['MTRANS']=le.transform(df['MTRANS'])`  
 C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\288179486.py:26: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.  
 Try using `.loc[row_indexer,col_indexer] = value` instead


See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/using\\_indexers.html](https://pandas.pydata.org/pandas-docs/stable/using_indexers.html)  
`df['NObeyesdad']=le.transform(df['NObeyesdad'])`

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 2085 entries, 0 to 2110
Data columns (total 17 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Gender                                2085 non-null   int32
 1   Age                                   2085 non-null   float64
 2   Height                               2085 non-null   float64
 3   Weight                               2085 non-null   float64
 4   family_history_with_overweight       2085 non-null   int32
 5   FAVC                                 2085 non-null   int32
 6   FCVC                                 2085 non-null   float64
 7   NCP                                  2085 non-null   float64
 8   CAEC                                 2085 non-null   int32
 9   SMOKE                                2085 non-null   int32
10   CH2O                                 2085 non-null   float64
11   SCC                                  2085 non-null   int32
12   FAF                                  2085 non-null   float64
13   TUE                                  2085 non-null   float64
14   CALC                                 2085 non-null   int32
15   MTRANS                               2085 non-null   int32
16   NObeyesdad                           2085 non-null   int32
dtypes: float64(8), int32(9)
memory usage: 219.9 KB
```

`sns.heatmap(df.corr())`

↔ <Axes: >



## ✓ Training and Testing

```
X=df.drop('NObeyesdad',axis=1)
y=df['NObeyesdad']
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
```

```
print("X_train = ",X_train.shape)
print("X_test = ",X_test.shape)
print("y_train = ",y_train.shape)
print("y_test = ",y_test.shape)
```

↔ X\_train = (1668, 16)  
X\_test = (417, 16)



```
y_train = (1668,)
y_test = (417,)
```

## ✓ Model 1 Linear Regression

```
from sklearn.linear_model import LinearRegression
li=LinearRegression()
li
```



```
▼ LinearRegression ⓘ ?
LinearRegression()
```

```
li.fit(X_train,y_train)
y_pred=li.predict(X_test)
accuracy=li.score(X_test,y_test)
accuracy
```



```
0.2179840640752998
```

```
from sklearn.metrics import mean_absolute_error,mean_squared_error
```

```
# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_pred)
print("Mean Absolute Error (MAE):", mae)
```

```
# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error (MSE):", mse)
```

```
# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print("Root Mean Squared Error (RMSE):", rmse)
```



```
Mean Absolute Error (MAE): 1.4145600564157936
Mean Squared Error (MSE): 2.902726370515661
Root Mean Squared Error (RMSE): 1.703738938486663
```

## ✓ Model 2 Logistic Regression

```
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
lr
```



```
▼ LogisticRegression ⓘ ?
LogisticRegression()
```

```
lr.fit(X_train,y_train)
y_pred=lr.predict(X_test)
accuracy=lr.score(X_test,y_test)
accuracy
```

➞ c:\Python310\lib\site-packages\sklearn\linear\_model\\_logistic.py:469: ConvergenceWarning  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
0.7170263788968825
```



```
from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score
```

```
print(mean_squared_error(y_test,y_pred))
print(mean_absolute_error(y_test,y_pred))
print(r2_score(y_test,y_pred))
```

➞ 2.431654676258993  
0.709832134292566  
0.34489425981873123

## ✓ Model 3 SVM

```
from sklearn.svm import SVC
svm_classifier=SVC(probability=True)
svm_classifier
```

➞ SVC(probability=True)

```
from sklearn.metrics import accuracy_score
```

```
svm_classifier.fit(X_train,y_train)
y_pred=svm_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```


➞ 0.5563549160671463

```
#Confusion matrix
from sklearn.metrics import confusion_matrix
```


```
cm=confusion_matrix(y_test,y_pred)
print("SVM Confusion Matrix: \n",cm)
```

```
⇒ SVM Confusion Matrix:
[[48  7  0  0  0  0  0]
 [26 22  1  0  0  8  4]
 [ 0  0 21  6 12  0 27]
 [ 0  0  1 36 34  0  0]
 [ 0  0  0  0 57  0  0]
 [ 3 18  3  0  0 22  9]
 [ 0  2 15  0  0  9 26]]
```

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('SVM Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```




	precision	recall	f1-score	support
0	0.62	0.87	0.73	55
1	0.45	0.36	0.40	61
2	0.51	0.32	0.39	66
3	0.86	0.51	0.64	71
4	0.55	1.00	0.71	57
5	0.56	0.40	0.47	55
6	0.39	0.50	0.44	52
accuracy			0.56	417
macro avg	0.56	0.57	0.54	417
weighted avg	0.57	0.56	0.54	417

## ✓ Model 4 Decision Tree


```
from sklearn.tree import DecisionTreeClassifier
```

```
dt_classifier=DecisionTreeClassifier()
dt_classifier
```




▼ DecisionTreeClassifier ⓘ ?  
 DecisionTreeClassifier()

```
dt_classifier.fit(X_train,y_train)
y_pred=dt_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```

 0.920863309352518


```
#Confusion Matrix
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print("Decision Tree Confusion Matrix: \n",cm)
```

 Decision Tree Confusion Matrix:

```
[[51  4  0  0  0  0  0]
 [ 5 50  0  0  0  6  0]
 [ 0  0 63  1  0  0  2]
 [ 0  0  1 70  0  0  0]
 [ 0  0  0  0 57  0  0]
 [ 0  5  1  0  0 46  3]
 [ 0  0  1  0  0  4 47]]
```

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Decision Tree Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.91	0.93	0.92	55
1	0.85	0.82	0.83	61
2	0.95	0.95	0.95	66
3	0.99	0.99	0.99	71
4	1.00	1.00	1.00	57
5	0.82	0.84	0.83	55
6	0.90	0.90	0.90	52
accuracy			0.92	417
macro avg	0.92	0.92	0.92	417
weighted avg	0.92	0.92	0.92	417

## ✓ Model 5 Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_classifier=RandomForestClassifier()  
rf_classifier
```



```
▼ RandomForestClassifier ⓘ ?  
RandomForestClassifier()
```

```
rf_classifier.fit(X_train,y_train)  
y_pred=rf_classifier.predict(X_test)  
accuracy=accuracy_score(y_test,y_pred)  
accuracy
```



```
0.9400479616306955
```


```
#Confusion Matrix  
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)  
print("Random Forest Confusion Matrix: \n",cm)
```



```
Random Forest Confusion Matrix:  
[[50  5  0  0  0  0  0]  
 [ 1 58  0  0  0  2  0]  
 [ 0  0 62  0  0  0  4]  
 [ 0  0  0 71  0  0  0]  
 [ 0  0  0  0 57  0  0]  
 [ 0  4  0  0  0 50  1]  
 [ 0  5  0  0  0  3 44]]
```

```
#Plot Confusion Matrix  
plt.figure(figsize=(8,6))  
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.title('Random Forest Confusion Matrix')  
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.98	0.91	0.94	55
1	0.81	0.95	0.87	61
2	1.00	0.94	0.97	66
3	1.00	1.00	1.00	71
4	1.00	1.00	1.00	57
5	0.91	0.91	0.91	55
6	0.90	0.85	0.87	52
accuracy			0.94	417
macro avg	0.94	0.94	0.94	417
weighted avg	0.94	0.94	0.94	417

## ✓ Model 6 K Neighbors

```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn_classifier=KNeighborsClassifier()
```

```
knn_classifier
```



```
▼ KNeighborsClassifier ⓘ ?  
KNeighborsClassifier()
```

```
knn_classifier.fit(X_train,y_train)
```

```
y_pred=knn_classifier.predict(X_test)
```

```
accuracy=accuracy_score(y_test,y_pred)
```

```
accuracy
```



```
0.8633093525179856
```

```
#Confusion Matrix
```

```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
```

```
print("K Neighbors Confusion Matrix: \n",cm)
```



```
K Neighbors Confusion Matrix:
```

```
[[54  1  0  0  0  0  0]  
 [13 33  1  0  0  9  5]  
 [ 0  0 61  0  0  0  5]  
 [ 0  0  1 70  0  0  0]  
 [ 0  0  0  2 55  0  0]  
 [ 0  5  1  0  0 48  1]  
 [ 0  2  5  1  0  5 39]]
```

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
```

```
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
```

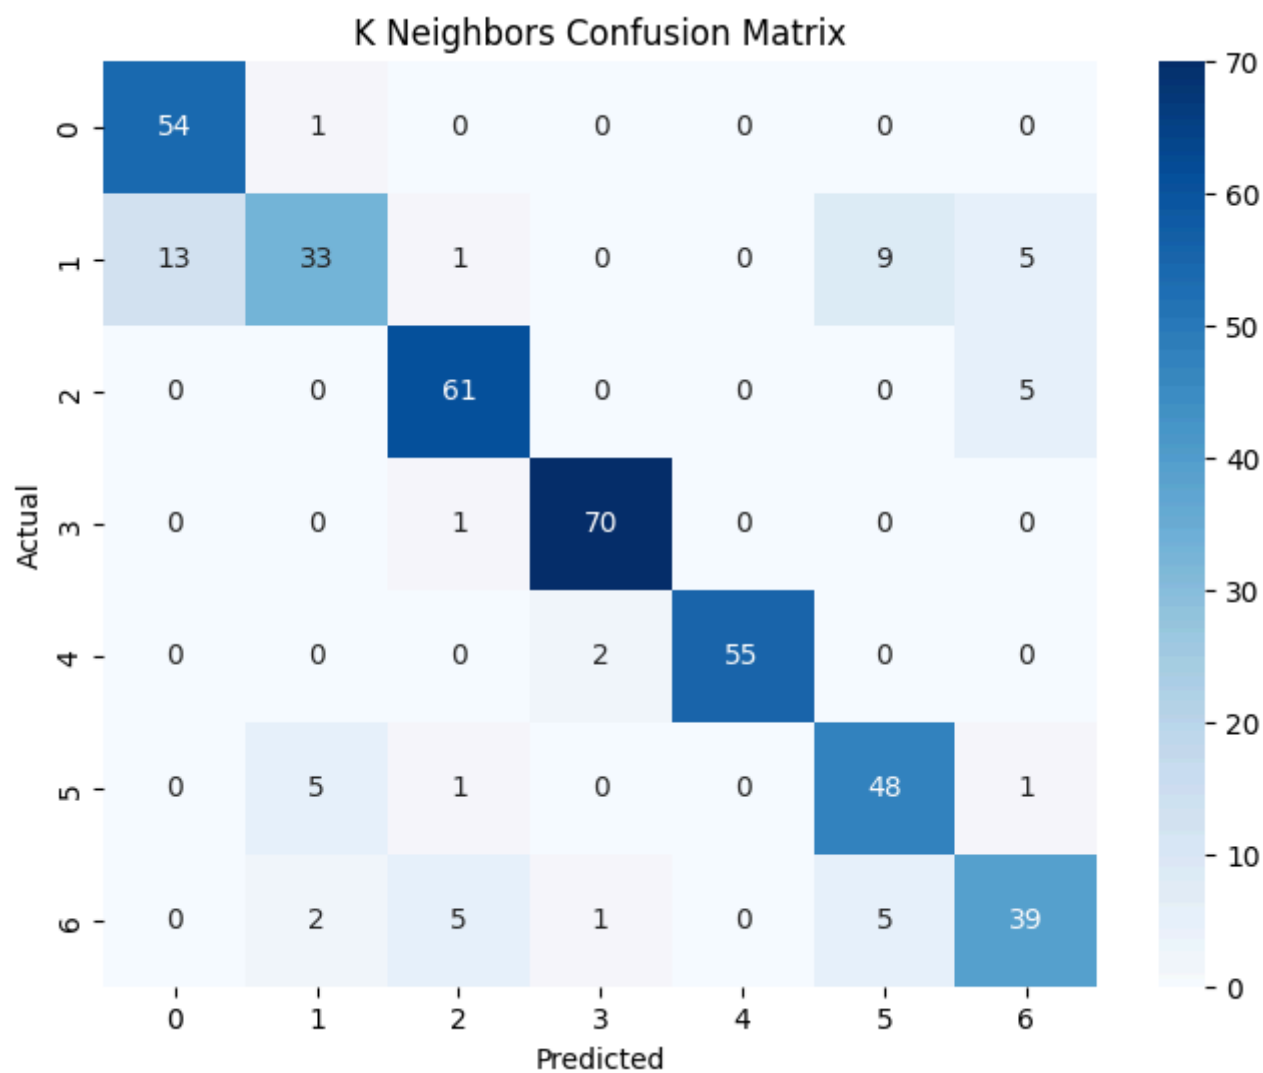
```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.title('K Neighbors Confusion Matrix')
```

```
plt.show()
```





```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.81	0.98	0.89	55
1	0.80	0.54	0.65	61
2	0.88	0.92	0.90	66
3	0.96	0.99	0.97	71
4	1.00	0.96	0.98	57
5	0.77	0.87	0.82	55
6	0.78	0.75	0.76	52
accuracy			0.86	417
macro avg	0.86	0.86	0.85	417
weighted avg	0.86	0.86	0.86	417

✓ Model 7 PLA

```
from sklearn.linear_model import Perceptron
```

```
pla_classifier=Perceptron(max_iter=1000,tol=1e-3,random_state=42)  
pla_classifier
```



▼ Perceptron ⓘ ?  
Perceptron(random\_state=42)

```
pla_classifier.fit(X_train,y_train)  
y_pred=pla_classifier.predict(X_test)  
accuracy=accuracy_score(y_test,y_pred)  
accuracy
```



0.4196642685851319

```
#Confusion Matrix  
from sklearn.metrics import confusion_matrix
```


```
cm=confusion_matrix(y_test,y_pred)  
print("PLA Matrix: \n",cm)
```



PLA Matrix:  
[[29 0 3 0 0 23 0]  
[ 8 1 22 1 0 26 3]  
[ 0 0 55 11 0 0 0]  
[ 0 0 11 60 0 0 0]  
[ 0 0 34 15 8 0 0]  
[ 0 0 27 4 0 22 2]  
[ 0 0 39 7 0 6 0]]

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))  
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.title('PLA Confusion Matrix')  
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.78	0.53	0.63	55
1	1.00	0.02	0.03	61
2	0.29	0.83	0.43	66
3	0.61	0.85	0.71	71
4	1.00	0.14	0.25	57
5	0.29	0.40	0.33	55
6	0.00	0.00	0.00	52
accuracy			0.42	417
macro avg	0.57	0.39	0.34	417
weighted avg	0.57	0.42	0.35	417

## ✓ Model 8 MLP

```
from sklearn.neural_network import MLPClassifier
```

```
mlp_classifier=MLPClassifier(hidden_layer_sizes=(100,),max_iter=300,random_state=42)
mlp_classifier
```



MLPClassifier

MLPClassifier(max\_iter=300, random\_state=42)

```
mlp_classifier.fit(X_train,y_train)
y_pred=mlp_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```



c:\Python310\lib\site-packages\sklearn\neural\_network\\_multilayer\_perceptron.py:691:  
warnings.warn(  
0.8225419664268585



#Confusion Matrix

```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print("MLP Matrix: \n",cm)
```




MLP Matrix:

```
[[47  8  0  0  0  0  0]
 [ 9 40  0  0  0  8  4]
 [ 0  0 56  3  0  1  6]
 [ 0  0  0 71  0  0  0]
 [ 0  0  0  0 57  0  0]
 [ 0  3  7  0  0 40  5]
 [ 0  2  6  1  0 11 32]]
```

#Plotting Confusion Matrix

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('MLP Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.84	0.85	0.85	55
1	0.75	0.66	0.70	61
2	0.81	0.85	0.83	66
3	0.95	1.00	0.97	71
4	1.00	1.00	1.00	57
5	0.67	0.73	0.70	55
6	0.68	0.62	0.65	52
accuracy			0.82	417
macro avg	0.81	0.81	0.81	417
weighted avg	0.82	0.82	0.82	417

## ✓ Model 9 Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
```

```
nb_classifier=GaussianNB()
```

```
nb_classifier
```



▼ GaussianNB ⓘ ?  
GaussianNB()

```
nb_classifier.fit(X_train,y_train)
```

```
y_pred=nb_classifier.predict(X_test)
```

```
accuracy=accuracy_score(y_test,y_pred)
```

```
accuracy
```



0.5827338129496403

```
#Confusion Matrix
```

```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
```

```
print("Naive Bayes Matrix: \n",cm)
```



Naive Bayes Matrix:

```
[[53  2  0  0  0  0  0]
 [44  6  5  0  0  3  3]
 [ 0  2 43 17  0  3  1]
 [ 0  0  2 68  0  0  1]
 [ 0  0  0  0 57  0  0]
 [11  2 30  1  0  8  3]
 [ 3  3 28  6  0  4  8]]
```

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
```


```
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.title('Naive Bayes Confusion Matrix')
```

```
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



	precision	recall	f1-score	support
0	0.48	0.96	0.64	55
1	0.40	0.10	0.16	61
2	0.40	0.65	0.49	66
3	0.74	0.96	0.83	71
4	1.00	1.00	1.00	57
5	0.44	0.15	0.22	55
6	0.50	0.15	0.24	52
accuracy			0.58	417
macro avg	0.57	0.57	0.51	417
weighted avg	0.57	0.58	0.52	417

## ✓ Model 10 KMeans Clustering

```
from sklearn.cluster import KMeans
```

```
kmeans=KMeans(n_clusters=3,random_state=42)
```

```
kmeans
```



KMeans



```
KMeans(n_clusters=3, random_state=42)
```

```
kmeans.fit(X)
```

```
y_pred=kmeans.predict(X)
```

```
from sklearn.metrics import silhouette_score
```

```
silhouette_avg=silhouette_score(X,y_pred)
```

```
print("silhouette_score: ",silhouette_avg)
```



```
silhouette_score: 0.5050222435923701
```

```
plt.figure(figsize=(8,6))
```

```
plt.scatter(X.iloc[:,0],X.iloc[:,1],c=y_pred,s=50,cmap='viridis')
```

```
centers=kmeans.cluster_centers_
```

```
plt.scatter(centers[:,0],centers[:,1],c='red',s=200,alpha=0.75)
```

```
plt.title('KMeans Clustering') #alpha=0.75 sets the transparency of the cluster cent
```

```
plt.xlabel('Feature 1')
```

```
plt.ylabel('Feature 2')
```

```
plt.show()
```





## KMeans Clustering

### ✓ Model 11 Dimensionality Reduction using PCA

```
from sklearn.decomposition import PCA
pca=PCA(n_components=2)
pca
```



PCA ⓘ ?  
PCA(n\_components=2)

```
X_pca=pca.fit_transform(X)
```

```
plt.figure(figsize=(8,6))
plt.scatter(X_pca[:,0],X_pca[:,1],cmap='viridis')
plt.title('Dimensionality Reduction using PCA')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```



C:\Users\rohit\AppData\Local\Temp\ipykernel\_16104\2795715780.py:2: UserWarning: No data points found in the specified range.  
plt.scatter(X\_pca[:,0],X\_pca[:,1],cmap='viridis')

