

✓ Import Necessary Package

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn as sk
```

```
data=pd.read_csv('C:\Rohith\Backup\Desktop\SEM 6\Machine Learning Lab\Practices\Online Sh
```

```
data.head()
```



| | Administrative | Administrative_Duration | Informational | Informational_Duration | Pr |
|---|----------------|-------------------------|---------------|------------------------|----|
| 0 | 0 | 0.0 | 0 | 0.0 | |
| 1 | 0 | 0.0 | 0 | 0.0 | |
| 2 | 0 | 0.0 | 0 | 0.0 | |
| 3 | 0 | 0.0 | 0 | 0.0 | |
| 4 | 0 | 0.0 | 0 | 0.0 | |

```
data.describe()
```



| | Administrative | Administrative_Duration | Informational | Informational_Duration |
|-------|----------------|-------------------------|---------------|------------------------|
| count | 12330.000000 | 12330.000000 | 12330.000000 | 12330.000000 |
| mean | 2.315166 | 80.818611 | 0.503569 | 34.472398 |
| std | 3.321784 | 176.779107 | 1.270156 | 140.749294 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 50% | 1.000000 | 7.500000 | 0.000000 | 0.000000 |
| 75% | 4.000000 | 93.256250 | 0.000000 | 0.000000 |
| max | 27.000000 | 3398.750000 | 24.000000 | 2549.375000 |

```
data.shape
```



```
(12330, 18)
```

```
data.info()
```

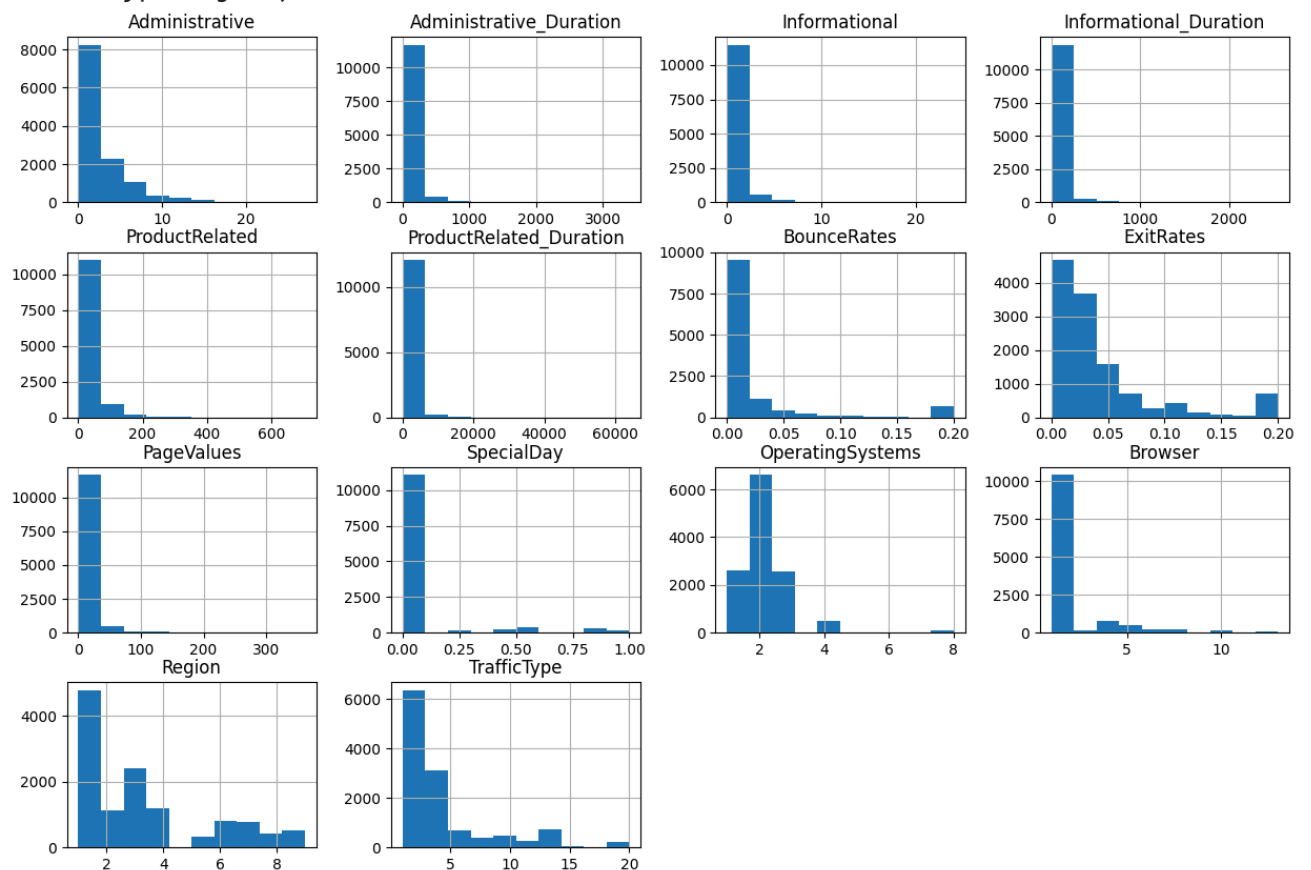
```
↩ <class 'pandas.core.frame.DataFrame'>  
RangeIndex: 12330 entries, 0 to 12329  
Data columns (total 18 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   Administrative                        12330 non-null  int64  
1   Administrative_Duration               12330 non-null  float64  
2   Informational                         12330 non-null  int64  
3   Informational_Duration               12330 non-null  float64  
4   ProductRelated                      12330 non-null  int64  
5   ProductRelated_Duration              12330 non-null  float64  
6   BounceRates                          12330 non-null  float64  
7   ExitRates                           12330 non-null  float64  
8   PageValues                          12330 non-null  float64  
9   SpecialDay                          12330 non-null  float64  
10  Month                               12330 non-null  object  
11  OperatingSystems                   12330 non-null  int64  
12  Browser                           12330 non-null  int64  
13  Region                            12330 non-null  int64  
14  TrafficType                       12330 non-null  int64  
15  VisitorType                       12330 non-null  object  
16  Weekend                           12330 non-null  bool  
17  Revenue                           12330 non-null  bool  
dtypes: bool(2), float64(7), int64(7), object(2)  
memory usage: 1.5+ MB
```

```
data.hist(figsize=(15,10))
```

```

array([[<Axes: title={'center': 'Administrative'}>,
       <Axes: title={'center': 'Administrative_Duration'}>,
       <Axes: title={'center': 'Informational'}>,
       <Axes: title={'center': 'Informational_Duration'}>],
 [<Axes: title={'center': 'ProductRelated'}>,
  <Axes: title={'center': 'ProductRelated_Duration'}>,
  <Axes: title={'center': 'BounceRates'}>,
  <Axes: title={'center': 'ExitRates'}>],
 [<Axes: title={'center': 'PageValues'}>,
  <Axes: title={'center': 'SpecialDay'}>,
  <Axes: title={'center': 'OperatingSystems'}>,
  <Axes: title={'center': 'Browser'}>],
 [<Axes: title={'center': 'Region'}>,
  <Axes: title={'center': 'TrafficType'}>],
 dtype=object)

```



```
#Checking the null values
data.isnull().sum()
```

```
⇒ Administrative      0
  Administrative_Duration  0
  Informational        0
  Informational_Duration  0
  ProductRelated       0
  ProductRelated_Duration  0
  BounceRates          0
  ExitRates            0
  PageValues           0
  SpecialDay           0
  Month               0
  OperatingSystems     0
  Browser              0
  Region              0
  TrafficType          0
  VisitorType          0
  Weekend              0
  Revenue              0
  dtype: int64
```

```
data.duplicated().sum()
```

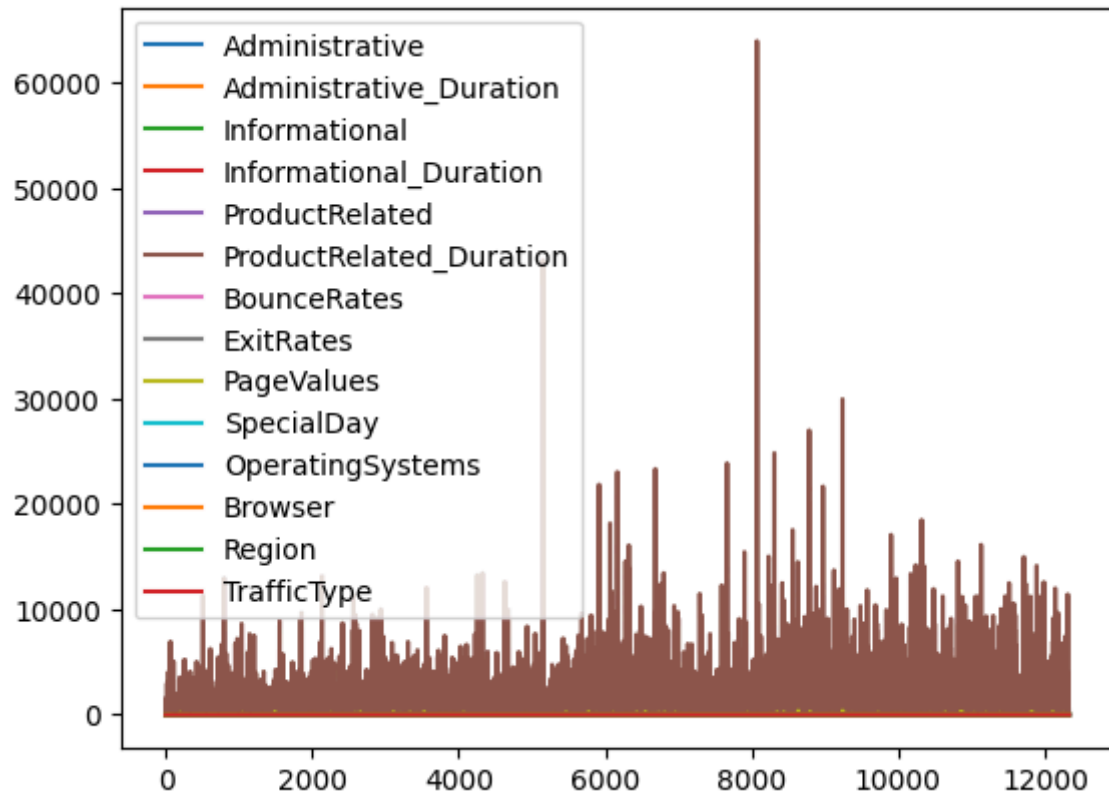
```
⇒ 125
```

```
data.drop_duplicates(inplace=True)
data.shape
```

```
⇒ (12205, 18)
```

```
data.plot()
```

↩ <Axes: >



```
data['ProductRelated_Duration'].unique()
```


↩ array([0. , 64. , 2.66666667, ..., 465.75 ,
 184.25 , 346.])

```
max(data['ProductRelated_Duration'].unique())
```

↩ 63973.52223

```
plt.boxplot(data['ProductRelated_Duration'])
```

```
➡ {'whiskers': [<matplotlib.lines.Line2D at 0x22a5e659270>,
<matplotlib.lines.Line2D at 0x22a5e659510>],
'caps': [<matplotlib.lines.Line2D at 0x22a5e6597b0>,
<matplotlib.lines.Line2D at 0x22a5e659a50>],
'boxes': [<matplotlib.lines.Line2D at 0x22a5efe2410>],
'medians': [<matplotlib.lines.Line2D at 0x22a5e659cf0>],
'fliers': [<matplotlib.lines.Line2D at 0x22a5e659ff0>],
'means': []}
```



```
df=data[data['ProductRelated_Duration']<=2500]
df
```




| | Administrative | Administrative_Duration | Informational | Informational_Duration |
|-------|----------------|-------------------------|---------------|------------------------|
| 0 | 0 | 0.0 | 0 | 0.0 |
| 1 | 0 | 0.0 | 0 | 0.0 |
| 2 | 0 | 0.0 | 0 | 0.0 |
| 3 | 0 | 0.0 | 0 | 0.0 |
| 4 | 0 | 0.0 | 0 | 0.0 |
| ... | ... | ... | ... | ... |
| 12325 | 3 | 145.0 | 0 | 0.0 |
| 12326 | 0 | 0.0 | 0 | 0.0 |
| 12327 | 0 | 0.0 | 0 | 0.0 |
| 12328 | 4 | 75.0 | 0 | 0.0 |
| 12329 | 0 | 0.0 | 0 | 0.0 |

10644 rows × 18 columns

df.plot()



<Axes: >




```
df['Revenue']=np.where(df['Revenue'],1,0)
plt.hist(df['Revenue'])
```

#Converting True=1 and False=0

➡ C:\Users\rohit\AppData\Local\Temp\ipykernel_17388\3012526398.py:1: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html#copy-on-write>

```
df['Revenue']=np.where(df['Revenue'],1,0)      #Converting True=1 and False=0
(array([9164.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        1480.]),
 array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ]),
 <BarContainer object of 10 artists>)
```




```
df['Weekend']=np.where(df['Weekend'],1,0)
plt.hist(df['Weekend'])
```


➡ C:\Users\rohit\AppData\Local\Temp\ipykernel_17388\1244464083.py:1: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/using_indexers.html

```
df['Weekend']=np.where(df['Weekend'],1,0)
(array([[8154.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        2490.]],
array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ]),
<BarContainer object of 10 artists>)
```




```
plt.scatter(x=df['SpecialDay'],y=df['TrafficType'])
```

 <matplotlib.collections.PathCollection at 0x22a5e41a860>



```
from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()          #Convert Object types into numeric values
```


```
le.fit(df['Month'])  
df['Month']=le.transform(df['Month'])
```

 C:\Users\rohit\AppData\Local\Temp\ipykernel_17388\3157372182.py:2: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html#copy-on-write>
df['Month']=le.transform(df['Month'])



```
le.fit(df['VisitorType'])  
df['VisitorType']=le.transform(df['VisitorType'])
```

 C:\Users\rohit\AppData\Local\Temp\ipykernel_17388\3298818332.py:2: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <https://pandas.pydata.org/pandas-docs/stable/10min.html#copy-on-write>
df['VisitorType']=le.transform(df['VisitorType'])




```
df.info()
```

```
➡ <class 'pandas.core.frame.DataFrame'>
Index: 10644 entries, 0 to 12329
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Administrative                        10644 non-null  int64
1   Administrative_Duration              10644 non-null  float64
2   Informational                        10644 non-null  int64
3   Informational_Duration               10644 non-null  float64
4   ProductRelated                      10644 non-null  int64
5   ProductRelated_Duration             10644 non-null  float64
6   BounceRates                         10644 non-null  float64
7   ExitRates                           10644 non-null  float64
8   PageValues                          10644 non-null  float64
9   SpecialDay                          10644 non-null  float64
10  Month                               10644 non-null  int32
11  OperatingSystems                    10644 non-null  int64
12  Browser                             10644 non-null  int64
13  Region                             10644 non-null  int64
14  TrafficType                         10644 non-null  int64
15  VisitorType                         10644 non-null  int32
16  Weekend                             10644 non-null  int32
17  Revenue                             10644 non-null  int32
dtypes: float64(7), int32(4), int64(7)
memory usage: 1.4 MB
```

```
sns.heatmap(df.corr())
```

↔ <Axes: >



✓ Training and testing

```
X=df.drop('Revenue',axis=1)
y=df['Revenue']
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
```

```
print("X_train = ",X_train.shape)
print("X_test = ",X_test.shape)
print("y_train = ",y_train.shape)
print("y_test = ",y_test.shape)
```

↔

```
X_train = (8515, 17)
X_test = (2129, 17)
y_train = (8515,)
y_test = (2129,)
```

✓ Model 1 Linear Regression

```
from sklearn.linear_model import LinearRegression
li=LinearRegression()
li
```



```
LinearRegression ⓘ ?
LinearRegression()
```

```
li.fit(X_train,y_train)
y_pred=li.predict(X_test)
accuracy=li.score(X_test,y_test)
accuracy
```



```
0.3305981754646069
```

```
from sklearn.metrics import mean_absolute_error,mean_squared_error
```

```
# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_pred)
print("Mean Absolute Error (MAE):", mae)
```

```
# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error (MSE):", mse)
```

```
# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print("Root Mean Squared Error (RMSE):", rmse)
```




```
Mean Absolute Error (MAE): 0.1747474516040395
Mean Squared Error (MSE): 0.0814856359227879
Root Mean Squared Error (RMSE): 0.28545688977985434
```

```
#Plotting ROC curve
from sklearn.metrics import roc_curve, auc, accuracy_score
```

```
fpr, tpr, threshold = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)
```

```
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend(loc="lower right")
plt.show()
```

➡ No artists with labels found to put in legend. Note that artists whose label start w



```

from sklearn.metrics import confusion_matrix

# Convert regression output to binary classification
threshold = y_test.mean() # Example threshold, you can set your own
y_test_class = (y_test > threshold).astype(int)
y_pred_class = (y_pred > threshold).astype(int)

# Compute confusion matrix
cm = confusion_matrix(y_test_class, y_pred_class)
accuracy = accuracy_score(y_test_class, y_pred_class)
print("Linear Regression Accuracy: \n ",accuracy)
print("Linear Regression Confusion Matrix: \n",cm)

```

```


➡ Linear Regression Accuracy:
  0.8271488961953969
Linear Regression Confusion Matrix:
[[1503  324]
 [  44 258]]

```

```

# Plot confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```



```
from sklearn.metrics import classification_report
print(classification_report(y_test_class,y_pred_class))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.97 | 0.82 | 0.89 | 1827 |
| 1 | 0.44 | 0.85 | 0.58 | 302 |
| accuracy | | | 0.83 | 2129 |
| macro avg | 0.71 | 0.84 | 0.74 | 2129 |
| weighted avg | 0.90 | 0.83 | 0.85 | 2129 |

✓ Model 2 Logistic Regression

```
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
lr
```




▼ LogisticRegression ⓘ ?
LogisticRegression()

```
lr.fit(X_train,y_train)
y_pred=lr.predict(X_test)
accuracy=lr.score(X_test,y_test)
accuracy
```



c:\Python310\lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarn
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
0.9060591827148896
```




```
#Plotting ROC curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color="darkorange",lw=2)
plt.plot([0,1],[0,1],color="navy",lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w




```
#Confusion matrix
from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred)
print("Logistic REgression Confusion matrix: \n",cm)
```

⇒ Logistic REgression Confusion matrix:

```
[[1790  37]
 [ 163 139]]
```

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.title('Logistic Regression Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.92 | 0.98 | 0.95 | 1827 |
| 1 | 0.79 | 0.46 | 0.58 | 302 |
| accuracy | | | 0.91 | 2129 |
| macro avg | 0.85 | 0.72 | 0.76 | 2129 |
| weighted avg | 0.90 | 0.91 | 0.90 | 2129 |

✓ Model 3 SVM

```
from sklearn.svm import SVC
svm_classifier=SVC(probability=True)
svm_classifier
```



▼ SVC ⓘ ?
SVC(probability=True)

```
svm_classifier.fit(X_train,y_train)
y_pred=svm_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```




0.8891498356035697

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1], [0,1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w




```
#Confusion matrix
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print("SVM Confusion Matrix: \n",cm)
```

⇒ SVM Confusion Matrix:

```
[[1812  15]
 [ 221  81]]
```

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('SVM Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.89 | 0.99 | 0.94 | 1827 |
| 1 | 0.84 | 0.27 | 0.41 | 302 |
| accuracy | | | 0.89 | 2129 |
| macro avg | 0.87 | 0.63 | 0.67 | 2129 |
| weighted avg | 0.88 | 0.89 | 0.86 | 2129 |

✓ Model 4 Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
```

```
dt_classifier=DecisionTreeClassifier()
dt_classifier
```



▼ DecisionTreeClassifier ⓘ ?

DecisionTreeClassifier()

```
dt_classifier.fit(X_train,y_train)
y_pred=dt_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```




0.8891498356035697

```
#Plotting ROC Curve
from sklearn.metrics import roc_curve,auc,accuracy_score

fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)

plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1], [0,1], color='navy',linestyle='--',lw=2)
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w




```
#Confusion Matrix
from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred)
print("Decision Tree Confusion Matrix: \n",cm)
```

⇒ Decision Tree Confusion Matrix:

```
[[1701  126]
 [ 110  192]]
```

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Decision Tree Confusion Matrix')
plt.show()
```

```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.94 | 0.93 | 0.94 | 1827 |
| 1 | 0.60 | 0.64 | 0.62 | 302 |
| accuracy | | | 0.89 | 2129 |
| macro avg | 0.77 | 0.78 | 0.78 | 2129 |
| weighted avg | 0.89 | 0.89 | 0.89 | 2129 |

✓ Model 5 Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_classifier=RandomForestClassifier()
rf_classifier
```



▼ RandomForestClassifier ⓘ ?

RandomForestClassifier()

```
rf_classifier.fit(X_train,y_train)
y_pred=rf_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```



0.9248473461719117


#Plotting ROC curve

```
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1], [0,1], color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w




```
#Confusion Matrix
from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred)
print("Random Forest Confusion Matrix: \n",cm)
```

⇒ Random Forest Confusion Matrix:

```
[[1776  51]
 [ 109 193]]
```

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Random Forest Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report  
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.94 | 0.97 | 0.96 | 1827 |
| 1 | 0.79 | 0.64 | 0.71 | 302 |
| accuracy | | | 0.92 | 2129 |
| macro avg | 0.87 | 0.81 | 0.83 | 2129 |
| weighted avg | 0.92 | 0.92 | 0.92 | 2129 |

Model6 K Neighbors

```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn_classifier=KNeighborsClassifier()  
knn_classifier
```



▼ KNeighborsClassifier ⓘ ?

KNeighborsClassifier()

```
knn_classifier.fit(X_train,y_train)
y_pred=knn_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```




0.8900892437764208

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score

fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)

plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w



```
#Confusion Matrix
from sklearn.metrics import confusion_matrix

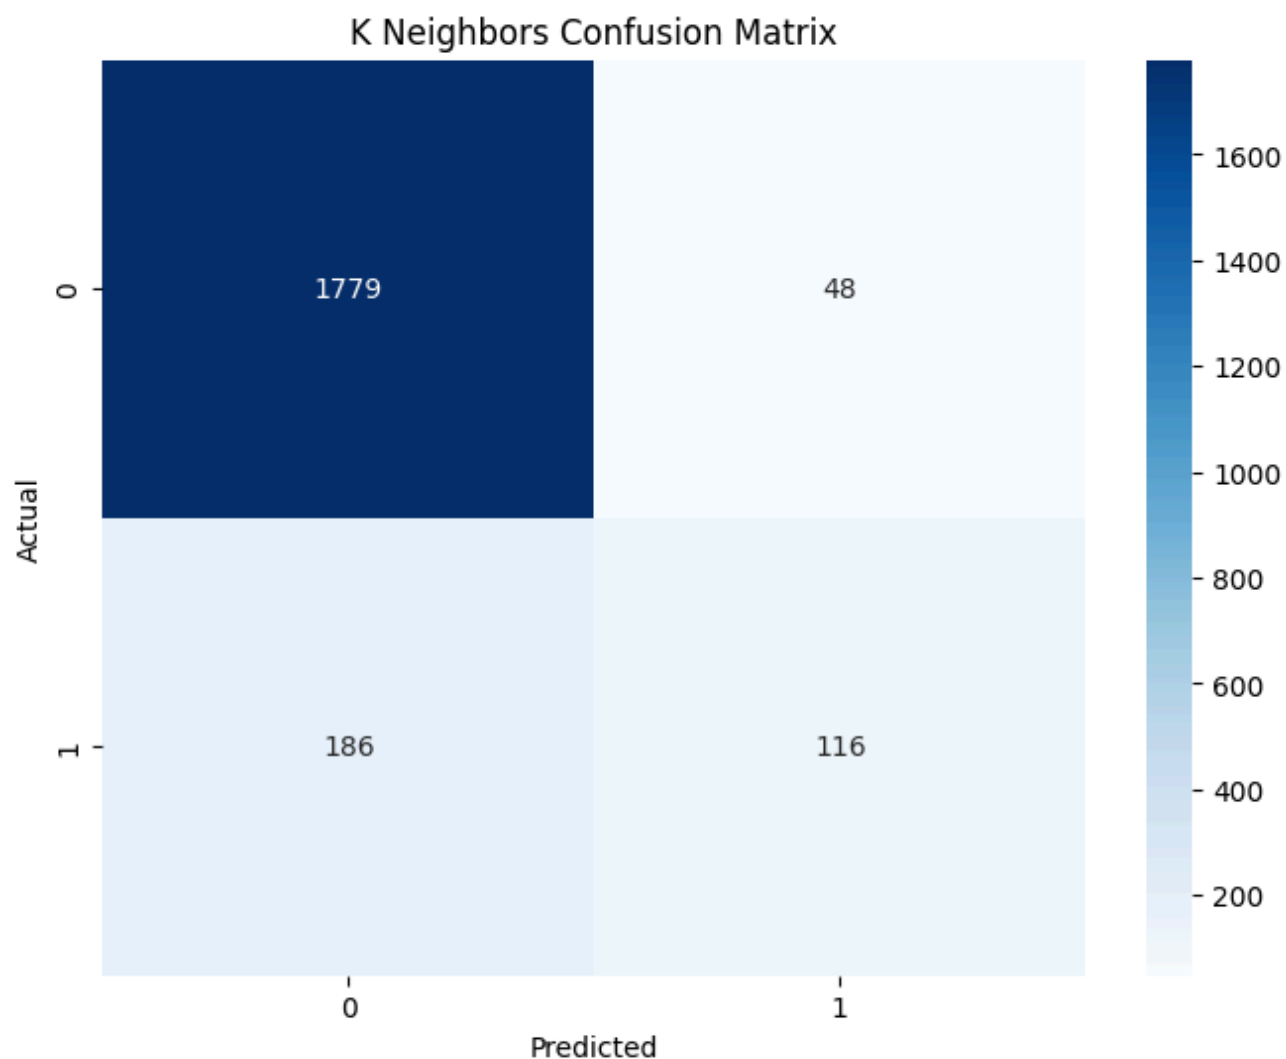
cm=confusion_matrix(y_test,y_pred)
print("K Neighbors Confusion Matrix: \n",cm)
```

⇒ K Neighbors Confusion Matrix:

```
[[1779  48]
 [ 186 116]]
```

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('K Neighbors Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.91 | 0.97 | 0.94 | 1827 |
| 1 | 0.71 | 0.38 | 0.50 | 302 |
| accuracy | | | 0.89 | 2129 |
| macro avg | 0.81 | 0.68 | 0.72 | 2129 |
| weighted avg | 0.88 | 0.89 | 0.88 | 2129 |

✓ Model 7 PLA Perceptron Learning Algorithm

```
from sklearn.linear_model import Perceptron
```

```
pla_classifier=Perceptron(max_iter=1000,tol=1e-3,random_state=42)
pla_classifier
```



▼ Perceptron ⓘ ?
Perceptron(random_state=42)

```
pla_classifier.fit(X_train,y_train)
y_pred=pla_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```




0.9013621418506341

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score

fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)

plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```


⇒ No artists with labels found to put in legend. Note that artists whose label start w



```
#Confusion Matrix
from sklearn.metrics import confusion_matrix
```

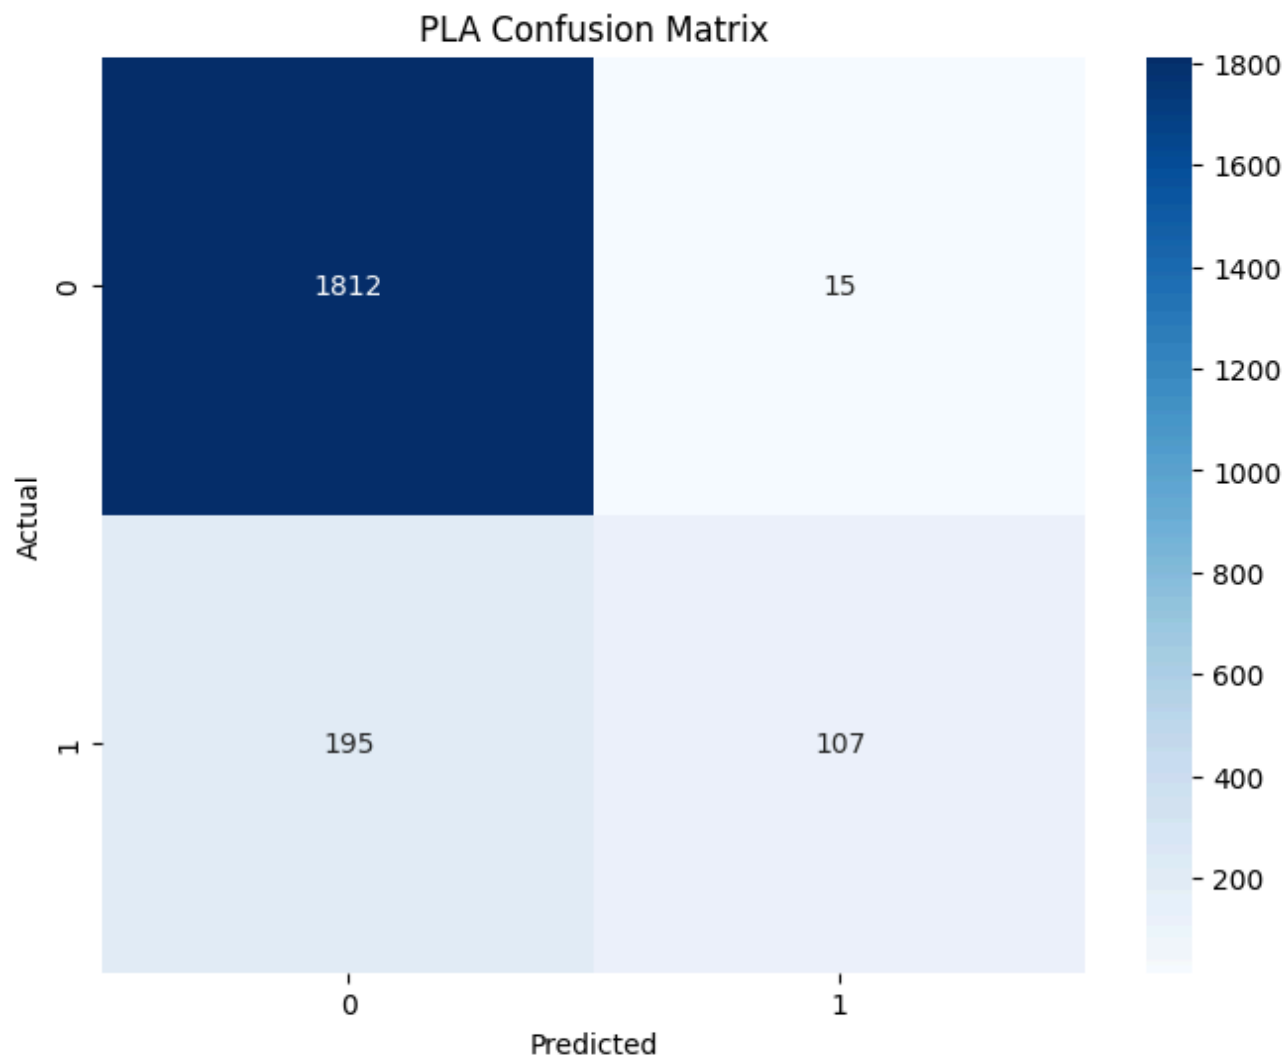
```
cm=confusion_matrix(y_test,y_pred)
print("PLA Matrix: \n",cm)
```

⇒ PLA Matrix:

```
[[1812  15]
 [ 195 107]]
```

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('PLA Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.90 | 0.99 | 0.95 | 1827 |
| 1 | 0.88 | 0.35 | 0.50 | 302 |
| accuracy | | | 0.90 | 2129 |
| macro avg | 0.89 | 0.67 | 0.72 | 2129 |
| weighted avg | 0.90 | 0.90 | 0.88 | 2129 |

✓ Model 8 MLP MultiLayer Perceptron

```
from sklearn.neural_network import MLPClassifier
```

```
mlp_classifier=MLPClassifier(hidden_layer_sizes=(100,),max_iter=300,random_state=42)
mlp_classifier
```



MLPClassifier



```
mlp_classifier.fit(X_train,y_train)
y_pred=mlp_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```



0.9173320807891029

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score

fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)

plt.figure()
```