

✓ Import necessary Package

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn as sk
import seaborn as sns
```

<https://archive.ics.uci.edu/dataset/850/raisin>

```
data=pd.read_excel('C:\Rohith\Backup\Desktop\SEM 6\Machine Learning Lab\Practices\Raisin\
```

```
data.head()
```



| | Area | MajorAxisLength | MinorAxisLength | Eccentricity | ConvexArea | Extent | Perim |
|---|-------|-----------------|-----------------|--------------|------------|----------|-------|
| 0 | 87524 | 442.246011 | 253.291155 | 0.819738 | 90546 | 0.758651 | 1184 |
| 1 | 75166 | 406.690687 | 243.032436 | 0.801805 | 78789 | 0.684130 | 1121 |
| 2 | 90856 | 442.267048 | 266.328318 | 0.798354 | 93717 | 0.637613 | 1208 |
| 3 | 45928 | 286.540559 | 208.760042 | 0.684989 | 47336 | 0.699599 | 844 |

```
data.shape
```



```
(900, 8)
```

```
data.describe()
```



| | Area | MajorAxisLength | MinorAxisLength | Eccentricity | ConvexArea | |
|-------|---------------|-----------------|-----------------|--------------|---------------|---|
| count | 900.000000 | 900.000000 | 900.000000 | 900.000000 | 900.000000 | 9 |
| mean | 87804.127778 | 430.929950 | 254.488133 | 0.781542 | 91186.090000 | |
| std | 39002.111390 | 116.035121 | 49.988902 | 0.090318 | 40769.290132 | |
| min | 25387.000000 | 225.629541 | 143.710872 | 0.348730 | 26139.000000 | |
| 25% | 59348.000000 | 345.442898 | 219.111126 | 0.741766 | 61513.250000 | |
| 50% | 78902.000000 | 407.803951 | 247.848409 | 0.798846 | 81651.000000 | |
| 75% | 105028.250000 | 494.187014 | 279.888575 | 0.842571 | 108375.750000 | |

```
data.info()
```








```
↩ <class 'pandas.core.frame.DataFrame'>  
RangeIndex: 900 entries, 0 to 899  
Data columns (total 8 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Area                  900 non-null   int64  
1   MajorAxisLength       900 non-null   float64  
2   MinorAxisLength       900 non-null   float64  
3   Eccentricity          900 non-null   float64  
4   ConvexArea            900 non-null   int64  
5   Extent                900 non-null   float64  
6   Perimeter             900 non-null   float64  
7   Class                 900 non-null   object  
dtypes: float64(5), int64(2), object(1)  
memory usage: 56.4+ KB
```

```
data.hist(figsize=(15,10))
```

```

array([[<Axes: title={'center': 'Area'}>,
       <Axes: title={'center': 'MajorAxisLength'}>,
       <Axes: title={'center': 'MinorAxisLength'}>],
      [<Axes: title={'center': 'Eccentricity'}>,
       <Axes: title={'center': 'ConvexArea'}>,
       <Axes: title={'center': 'Extent'}>],
      [<Axes: title={'center': 'Perimeter'}>, <Axes: >, <Axes: >]],
      dtype=object)

```



```
data.isnull().sum()
```

```
⇒ Area          0
MajorAxisLength  0
MinorAxisLength  0
Eccentricity     0
ConvexArea       0
Extent          0
Perimeter        0
Class           0
dtype: int64
```

```
data.dropna(inplace=True)
```

```
data.duplicated().sum()
```

```
⇒ 0
```

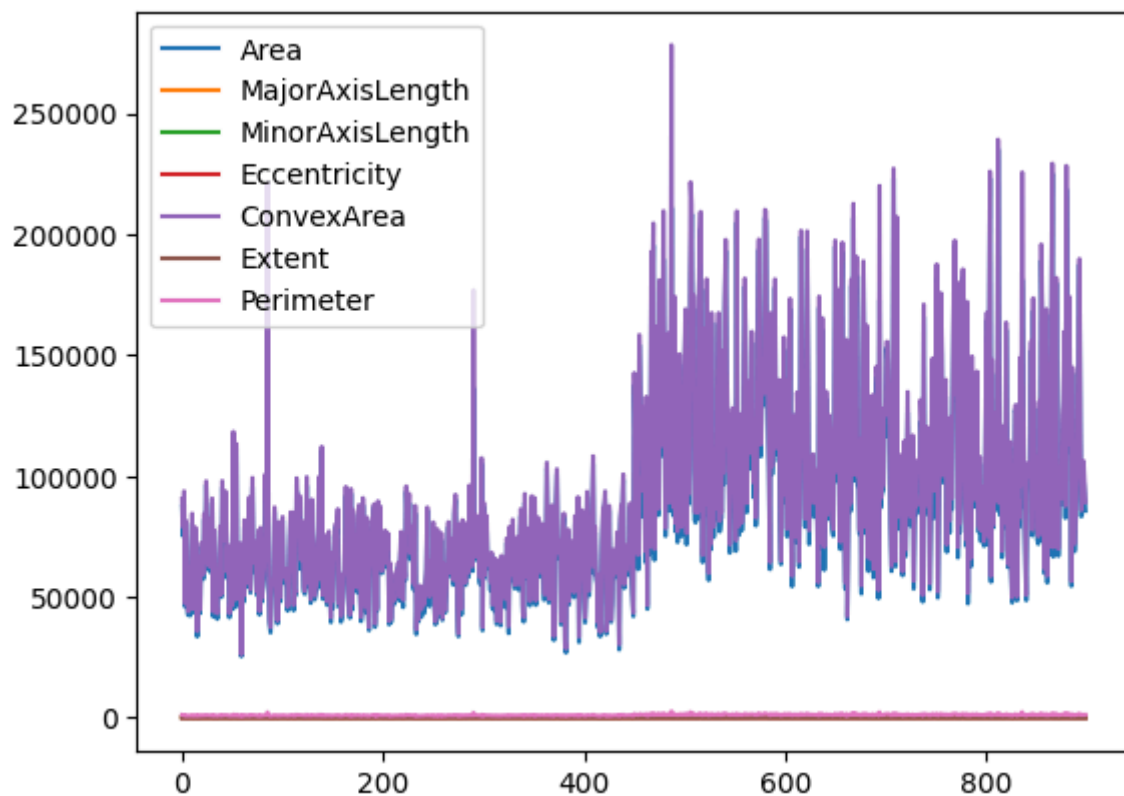
```
data.drop_duplicates(inplace=True)
```

```
data.shape
```

```
⇒ (900, 8)
```

```
data.plot()
```

```
⇒ <Axes: >
```




```
max(data['ConvexArea'].unique())
```

```
⇒ 278217
```

```
plt.boxplot(data['ConvexArea'])
```

```
⇒ { 'whiskers': [<matplotlib.lines.Line2D at 0x13a1667e920>,
<matplotlib.lines.Line2D at 0x13a1667ebc0>],
'caps': [<matplotlib.lines.Line2D at 0x13a1667ee60>,
<matplotlib.lines.Line2D at 0x13a1667f100>],
'boxes': [<matplotlib.lines.Line2D at 0x13a1667e680>],
'medians': [<matplotlib.lines.Line2D at 0x13a1667f3a0>],
'fliers': [<matplotlib.lines.Line2D at 0x13a1667f640>],
'means': []}
```




```
df=data[data['ConvexArea']<=240000]
```

```
df.shape
```

```
⇒ (899, 8)
```

```
df.plot()
```



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 899 entries, 0 to 899
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Area                   899 non-null    int64
1   MajorAxisLength        899 non-null    float64
2   MinorAxisLength        899 non-null    float64
3   Eccentricity            899 non-null    float64
4   ConvexArea              899 non-null    int64
5   Extent                  899 non-null    float64
6   Perimeter               899 non-null    float64
7   Class                   899 non-null    object
dtypes: float64(5), int64(2), object(1)
memory usage: 63.2+ KB
```

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
le
```

```
LabelEncoder
LabelEncoder()
```

```
le.fit(df['Class'])
df['Class']=le.transform(df['Class'])
```

➡ C:\Users\rohit\AppData\Local\Temp\ipykernel_2976\1694167980.py:2: SettingWithCopyWarning
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/using_indexers.html
df['Class']=le.transform(df['Class'])




df.info()

➡ <class 'pandas.core.frame.DataFrame'>
Index: 899 entries, 0 to 899
Data columns (total 8 columns):
Column Non-Null Count Dtype
--- -
0 Area 899 non-null int64
1 MajorAxisLength 899 non-null float64
2 MinorAxisLength 899 non-null float64
3 Eccentricity 899 non-null float64
4 ConvexArea 899 non-null int64
5 Extent 899 non-null float64
6 Perimeter 899 non-null float64
7 Class 899 non-null int32
dtypes: float64(5), int32(1), int64(2)
memory usage: 59.7 KB

sns.heatmap(df.corr())

↔ <Axes: >



✓ Training and Testing Data

```
X=df.drop('Class',axis=1)
y=df['Class']
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
```

```
print("X_train = ",X_train.shape)
print("X_test = ",X_test.shape)
print("y_train = ",y_train.shape)
print("y_test = ",y_test.shape)
```

↔

```
X_train = (719, 7)
X_test = (180, 7)
y_train = (719,)
y_test = (180,)
```


✓ Model 1 Linear Regression

```
from sklearn.linear_model import LinearRegression
li=LinearRegression()
li
```



▼ LinearRegression ⓘ ?
LinearRegression()

```
li.fit(X_train,y_train)
y_pred=li.predict(X_test)
accuracy=li.score(X_test,y_test)
accuracy
```



0.42138975251373

```
from sklearn.metrics import mean_absolute_error,mean_squared_error
```

```
from sklearn.metrics import mean_absolute_error,mean_squared_error
```

```
# Calculate Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_pred)
print("Mean Absolute Error (MAE):", mae)
```

```
# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error (MSE):", mse)
```

```
# Calculate Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print("Root Mean Squared Error (RMSE):", rmse)
```




Mean Absolute Error (MAE): 0.30527313252470945
Mean Squared Error (MSE): 0.14286672777438766
Root Mean Squared Error (RMSE): 0.3779771524502343

```
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)  
roc_auc=auc(fpr,tpr)
```

```
plt.figure()  
plt.plot(fpr,tpr,color='darkorange',lw=2)  
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')  
plt.xlim([0.0,1.0])  
plt.ylim([0.0,1.05])  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('ROC Curve')  
plt.legend(loc='lower right')  
plt.show()
```

➞ No artists with labels found to put in legend. Note that artists whose label start w



✓ Model 2 Logistic Regression

```
from sklearn.linear_model import LogisticRegression  
lr=LogisticRegression()  
lr
```

➞

▼ LogisticRegression ⓘ ?

LogisticRegression()

```
lr.fit(X_train,y_train)
y_pred=lr.predict(X_test)
accuracy=lr.score(X_test,y_test)
accuracy
```

➡ c:\Python310\lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
0.8111111111111111
```




```
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w




```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
print(cm)
```

⇒ $\begin{bmatrix} 63 & 17 \\ 17 & 83 \end{bmatrix}$

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```





| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.79 | 0.79 | 0.79 | 80 |
| 1 | 0.83 | 0.83 | 0.83 | 100 |
| accuracy | | | 0.81 | 180 |
| macro avg | 0.81 | 0.81 | 0.81 | 180 |
| weighted avg | 0.81 | 0.81 | 0.81 | 180 |

✓ Model 3 SVM

```
from sklearn.svm import SVC
svm_classifier=SVC(probability=True)
svm_classifier
```



SVC  
SVC(probability=True)

```
svm_classifier.fit(X_train,y_train)
y_pred=svm_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```



0.8111111111111111


```
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```



No artists with labels found to put in legend. Note that artists whose label start w




```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)  
print(cm)
```

```
↩ [[62 18]  
   [16 84]]
```

```
plt.figure(figsize=(8,6))  
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.title('Confusion Matrix')  
plt.show()
```



```
from sklearn.metrics import classification_report  
print(classification_report(y_test,y_pred))
```

```
↩
```

| | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| 0 | 0.79 | 0.78 | 0.78 | 80 |
| 1 | 0.82 | 0.84 | 0.83 | 100 |
| accuracy | | | 0.81 | 180 |

| | | | | |
|--------------|------|------|------|-----|
| macro avg | 0.81 | 0.81 | 0.81 | 180 |
| weighted avg | 0.81 | 0.81 | 0.81 | 180 |

✓ Model 4 Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
```

```
dt_classifier=DecisionTreeClassifier()  
dt_classifier
```



```
▼ DecisionTreeClassifier ⓘ ?  
DecisionTreeClassifier()
```

```
dt_classifier.fit(X_train,y_train)  
y_pred=dt_classifier.predict(X_test)  
accuracy=accuracy_score(y_test,y_pred)  
accuracy
```




```
0.75
```

```
#Plotting ROC Curve  
from sklearn.metrics import roc_curve,auc,accuracy_score  
  
fpr,tpr,threshold=roc_curve(y_test,y_pred)  
roc_auc=auc(fpr,tpr)
```

```
plt.figure()  
plt.plot(fpr,tpr,color='darkorange',lw=2)  
plt.plot([0,1], [0,1], color='navy',linestyle='--',lw=2)  
plt.xlim([0.0,1.0])  
plt.ylim([0.0,1.05])  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('ROC curve')  
plt.legend(loc='lower right')  
plt.show()
```


⇒ No artists with labels found to put in legend. Note that artists whose label start w




```
#Confusion Matrix
from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred)
print("Decision Tree Confusion Matrix: \n",cm)
```

⇒ Decision Tree Confusion Matrix:

```
[[64 16]
 [29 71]]
```

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Decision Tree Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.69 | 0.80 | 0.74 | 80 |
| 1 | 0.82 | 0.71 | 0.76 | 100 |
| accuracy | | | 0.75 | 180 |
| macro avg | 0.75 | 0.76 | 0.75 | 180 |
| weighted avg | 0.76 | 0.75 | 0.75 | 180 |

✓ Model 5 RandomForest

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_classifier=RandomForestClassifier()
rf_classifier
```



▼ RandomForestClassifier ⓘ ?

RandomForestClassifier()

```
rf_classifier.fit(X_train,y_train)
y_pred=rf_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```



0.8111111111111111


#Plotting ROC curve

```
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1], [0,1], color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w




```
#Confusion Matrix
from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test,y_pred)
print("Random Forest Confusion Matrix: \n",cm)
```

⇒ Random Forest Confusion Matrix:

```
[[64 16]
 [18 82]]
```

```
#Plot Confusion Matrix
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Random Forest Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report  
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.80 | 0.79 | 80 |
| 1 | 0.84 | 0.82 | 0.83 | 100 |
| accuracy | | | 0.81 | 180 |
| macro avg | 0.81 | 0.81 | 0.81 | 180 |
| weighted avg | 0.81 | 0.81 | 0.81 | 180 |

✓ Model 6 KNeighbors

```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn_classifier=KNeighborsClassifier()  
knn_classifier
```



▼ KNeighborsClassifier ⓘ ?

KNeighborsClassifier()

```
knn_classifier.fit(X_train,y_train)
y_pred=knn_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```




0.8055555555555556

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score

fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)

plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w



```
#Confusion Matrix
from sklearn.metrics import confusion_matrix

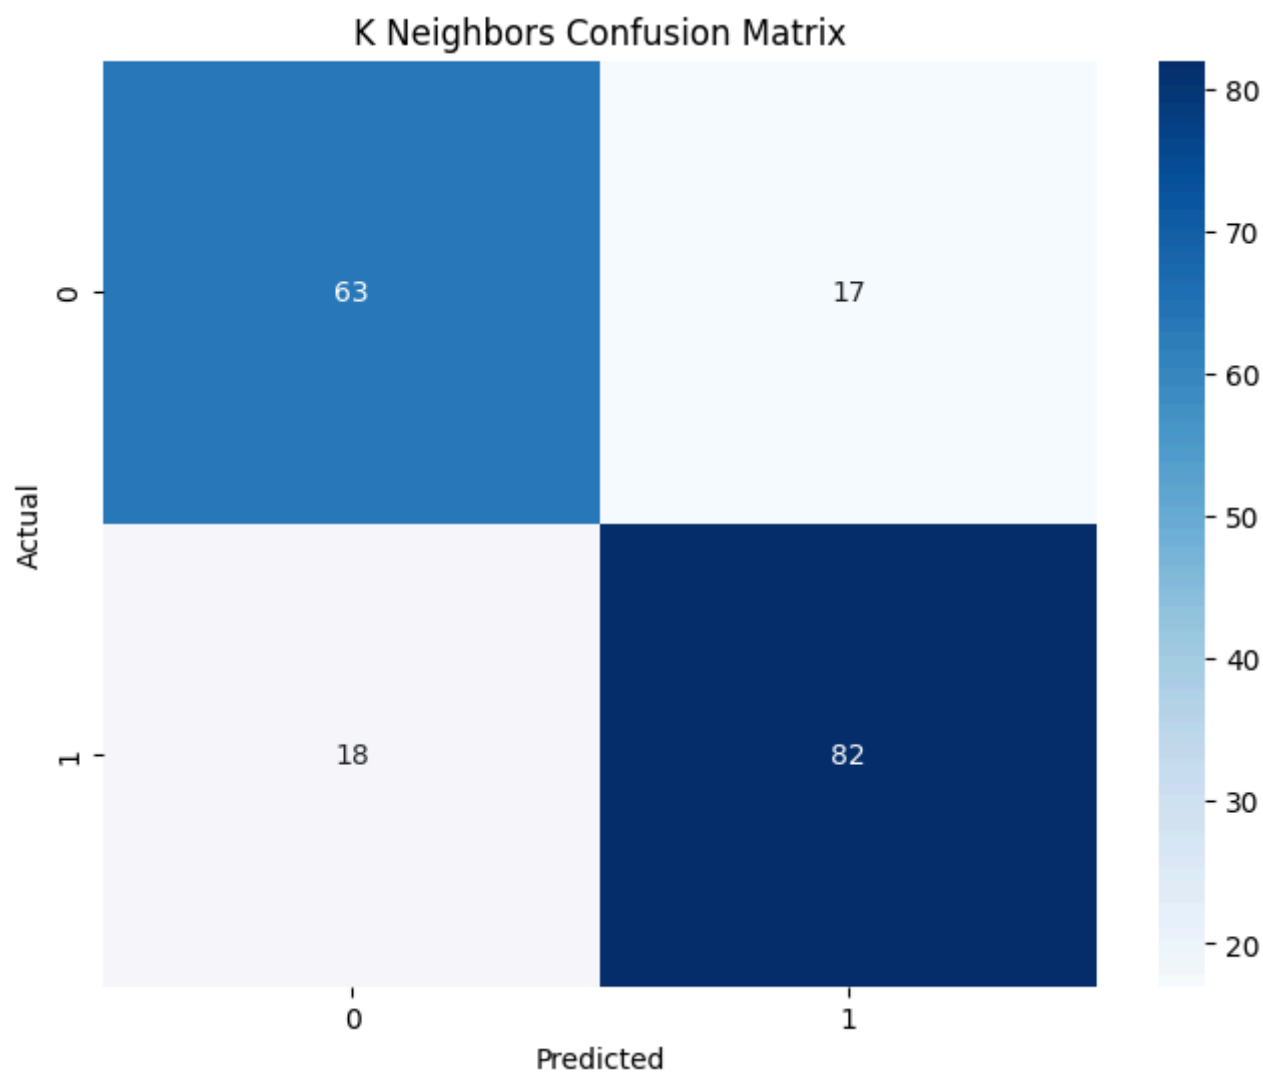
cm=confusion_matrix(y_test,y_pred)
print("K Neighbors Confusion Matrix: \n",cm)
```

⇒ K Neighbors Confusion Matrix:

```
[[63 17]
 [18 82]]
```

```
#Plotting Confusion Matrix

plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('K Neighbors Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.79 | 0.78 | 80 |
| 1 | 0.83 | 0.82 | 0.82 | 100 |
| accuracy | | | 0.81 | 180 |
| macro avg | 0.80 | 0.80 | 0.80 | 180 |
| weighted avg | 0.81 | 0.81 | 0.81 | 180 |

✓ Model 7 PLA

```
from sklearn.linear_model import Perceptron
```

```
pla_classifier=Perceptron(max_iter=1000,tol=1e-3,random_state=42)
pla_classifier
```




▼ Perceptron ⓘ ?
Perceptron(random_state=42)

```
pla_classifier.fit(X_train,y_train)
y_pred=pla_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```




0.4444444444444444

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score

fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)

plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w



```
#Confusion Matrix
from sklearn.metrics import confusion_matrix
```

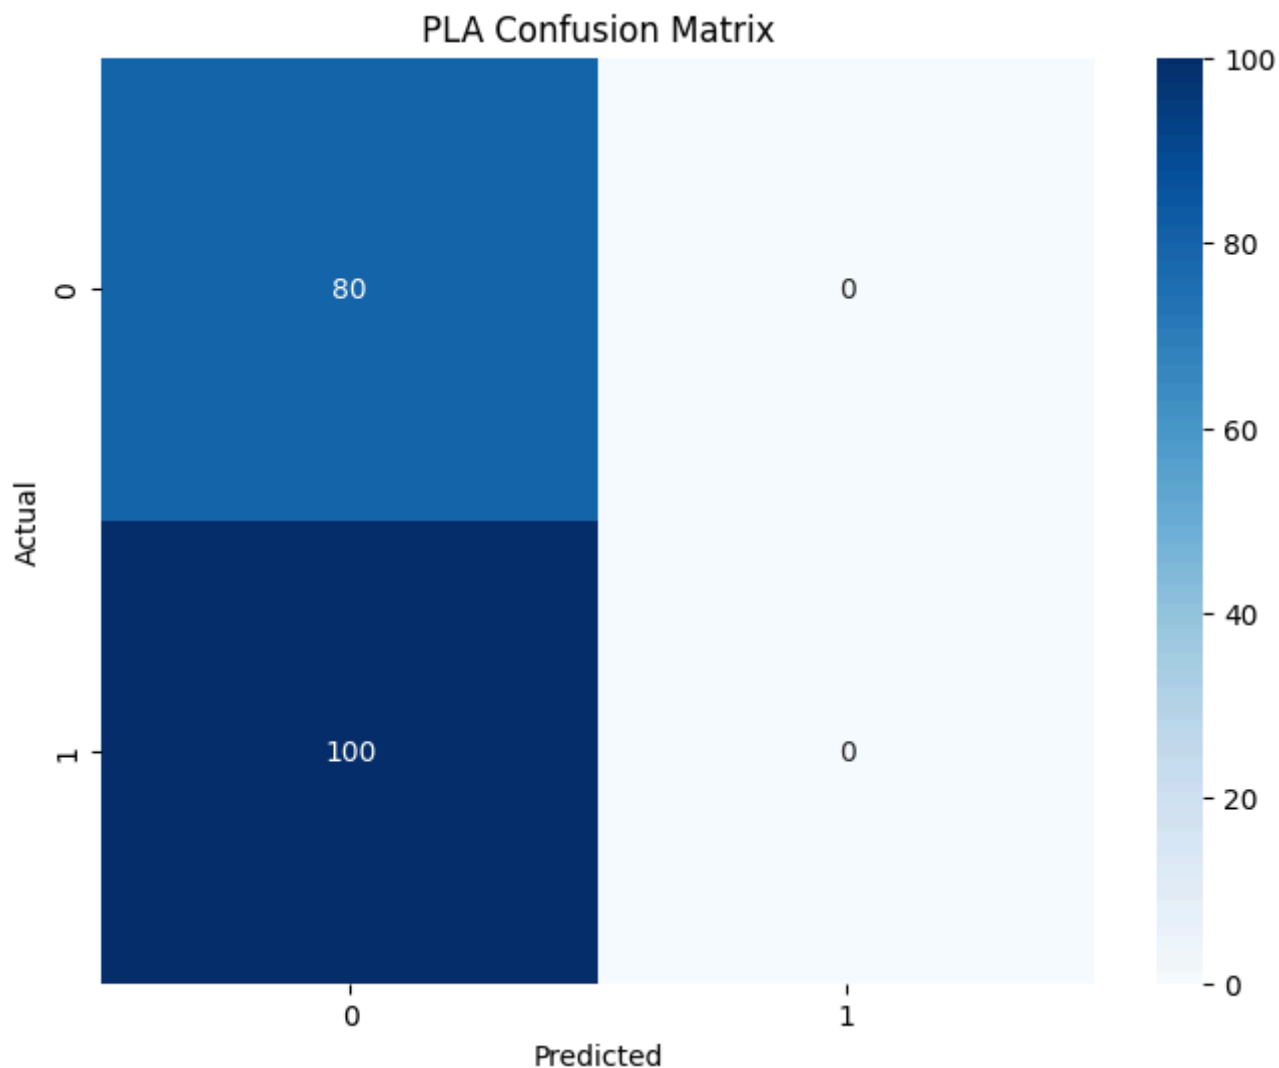
```
cm=confusion_matrix(y_test,y_pred)
print("PLA Matrix: \n",cm)
```

⇒ PLA Matrix:

```
[[ 80  0]
 [100  0]]
```

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('PLA Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.44 | 1.00 | 0.62 | 80 |
| 1 | 0.00 | 0.00 | 0.00 | 100 |
| accuracy | | | 0.44 | 180 |
| macro avg | 0.22 | 0.50 | 0.31 | 180 |
| weighted avg | 0.20 | 0.44 | 0.27 | 180 |

```
c:\Python310\lib\site-packages\sklearn\metrics\_classification.py:1509: UndefinedMetricWarning: Precision score is ill-defined: No samples predicted as positive
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Python310\lib\site-packages\sklearn\metrics\_classification.py:1509: UndefinedMetricWarning: Recall score is ill-defined: No samples predicted as positive
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Python310\lib\site-packages\sklearn\metrics\_classification.py:1509: UndefinedMetricWarning: F-score score is ill-defined: No samples predicted as positive
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```



✓ Model 8 MLP

```
from sklearn.neural_network import MLPClassifier
```

```
mlp_classifier=MLPClassifier(max_iter=300,hidden_layer_sizes=(100,),random_state=42)
mlp_classifier
```



MLPClassifier

MLPClassifier(max_iter=300, random_state=42)

```
mlp_classifier.fit(X_train,y_train)
y_pred=mlp_classifier.predict(X_test)
accuracy=accuracy_score(y_test,y_pred)
accuracy
```




0.8222222222222222

```
#Plotting ROC_curve
from sklearn.metrics import roc_curve,auc,accuracy_score
```

```
fpr,tpr,threshold=roc_curve(y_test,y_pred)
roc_auc=auc(fpr,tpr)
```

```
plt.figure()
plt.plot(fpr,tpr,color='darkorange',lw=2)
plt.plot([0,1],[0,1],color='navy',lw=2,linestyle='--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc='lower right')
plt.show()
```

⇒ No artists with labels found to put in legend. Note that artists whose label start w



```
#Confusion Matrix
from sklearn.metrics import confusion_matrix
```

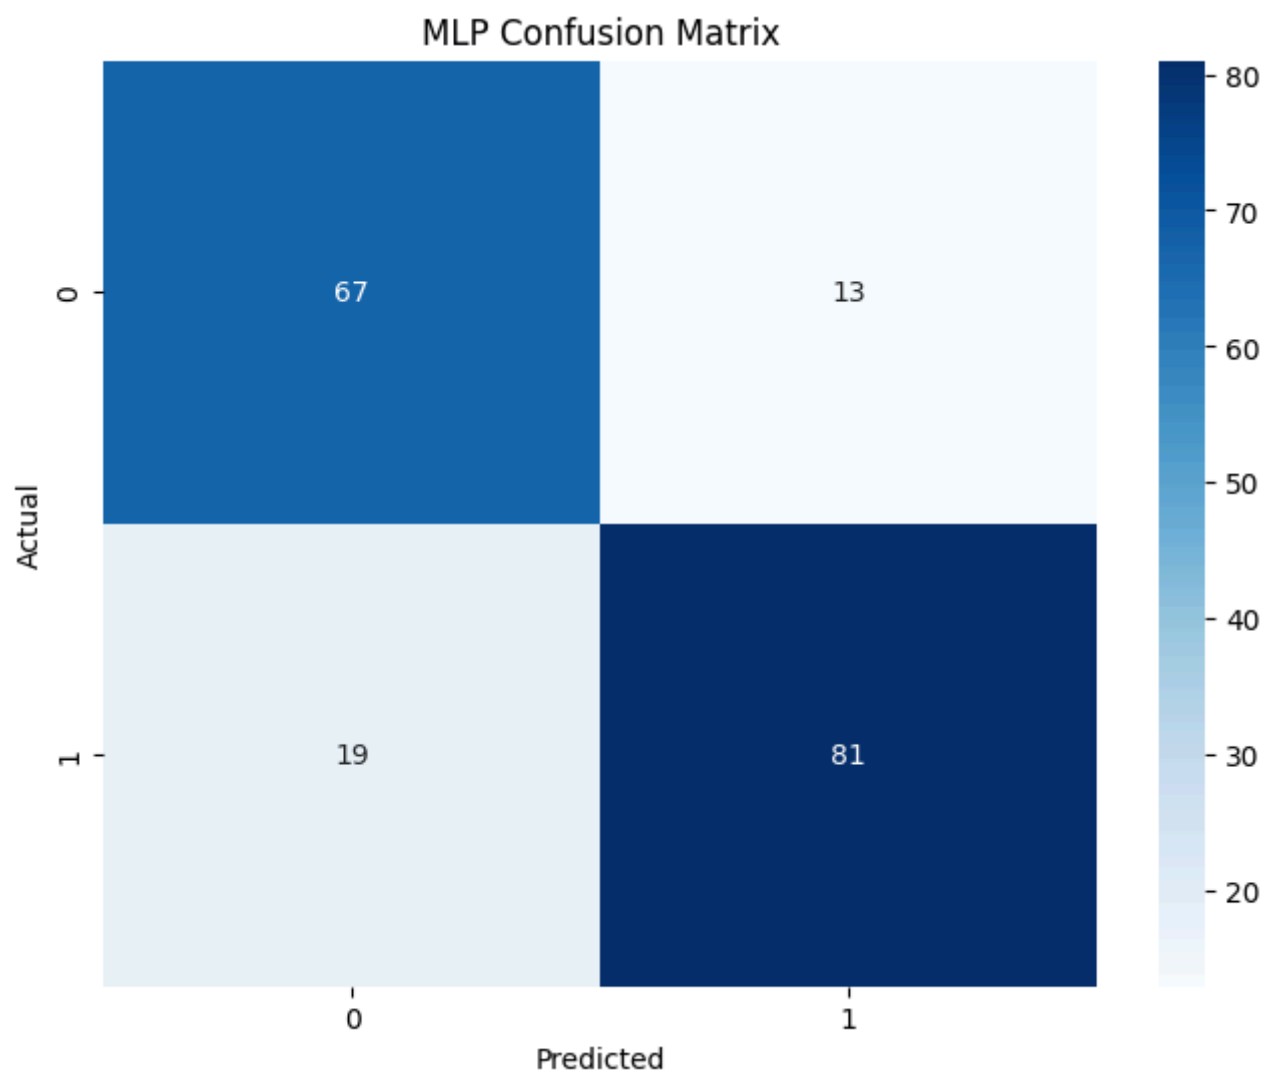
```
cm=confusion_matrix(y_test,y_pred)
print("MLP Matrix: \n",cm)
```

⇒ MLP Matrix:

```
[[67 13]
 [19 81]]
```

```
#Plotting Confusion Matrix
```

```
plt.figure(figsize=(8,6))
sns.heatmap(cm,annot=True,fmt='d',cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('MLP Confusion Matrix')
plt.show()
```



```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```



| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.84 | 0.81 | 80 |
| 1 | 0.86 | 0.81 | 0.84 | 100 |
| accuracy | | | 0.82 | 180 |
| macro avg | 0.82 | 0.82 | 0.82 | 180 |
| weighted avg | 0.82 | 0.82 | 0.82 | 180 |

✓ Model 9 Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
```

```
nb_classifier=GaussianNB()
nb_classifier
```



GaussianNB



```
nb_classifier.fit(X_train,y_train)
```